Introduction

Jun Zhu

dcszj@mail.tsinghua.edu.cn

http://ml.cs.tsinghua.edu.cn/~jun

Sate Key Lab of Intelligent Tech. & Systems,

Tsinghua University

Goals of this Lecture …

- Show that machine learning (ML) is cool
- Get you excited about ML
- Give an overview of basic problems & methods in ML
- Help you distinguish hype and science
- Entice you to take further study on ML, write a thesis on ML, dedicate your life to ML …

The age of Big Data

CERN Collider

320 x 10^{12} bytes/second

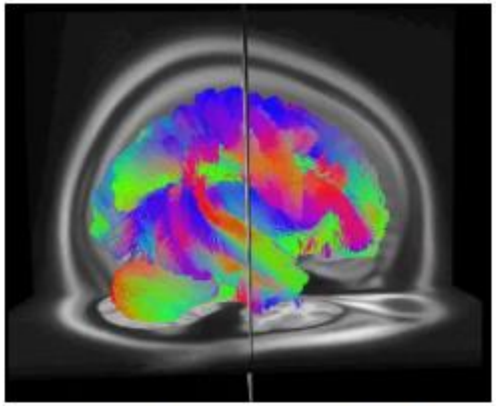

Prof. Tim Verstynen, CMU

Personal Connectome

10^{18} bytes/human

facebook

1 billion messages/day

twitter

200 million tweets/day

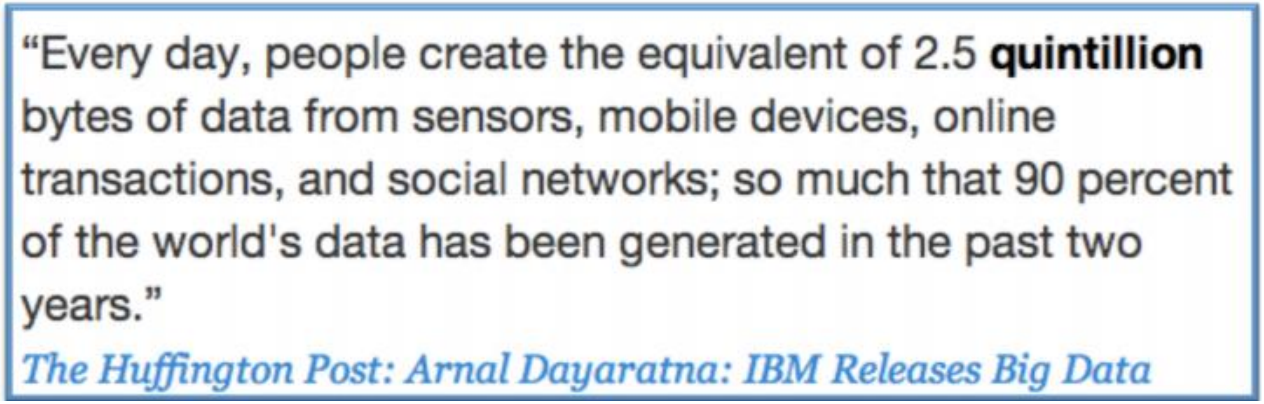

"Every day, people create the equivalent of 2.5 **quintillion** bytes of data from sensors, mobile devices, online transactions, and social networks; so much that 90 percent of the world's data has been generated in the past two years."

The Huffington Post: Arnal Dayaratna: IBM Releases Big Data

The age of Big Data

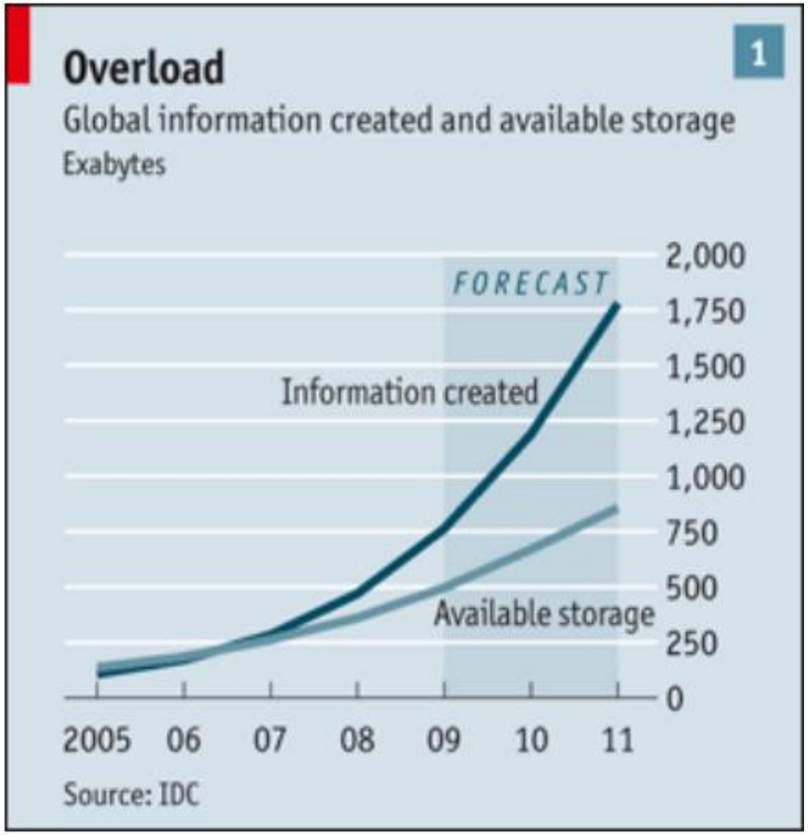

40,000 Exabytes by 2020
(IDC)

200 million in government funding
(White house initiative)

jobs shortage of 200,000 data experts by 2018
(Bloomberg)

"the sexiest job of the 21st century."
(Harvard Business Review)

Data ⇏ Knowledge

From Data to Knowledge …

What is Machine Learning?

Machine learning, a branch of artificial intelligence, is a scientific discipline concerned with the design and development of algorithms that take as input empirical data, and yield patterns or predictions thought to be features of the underlying mechanism that generated the data

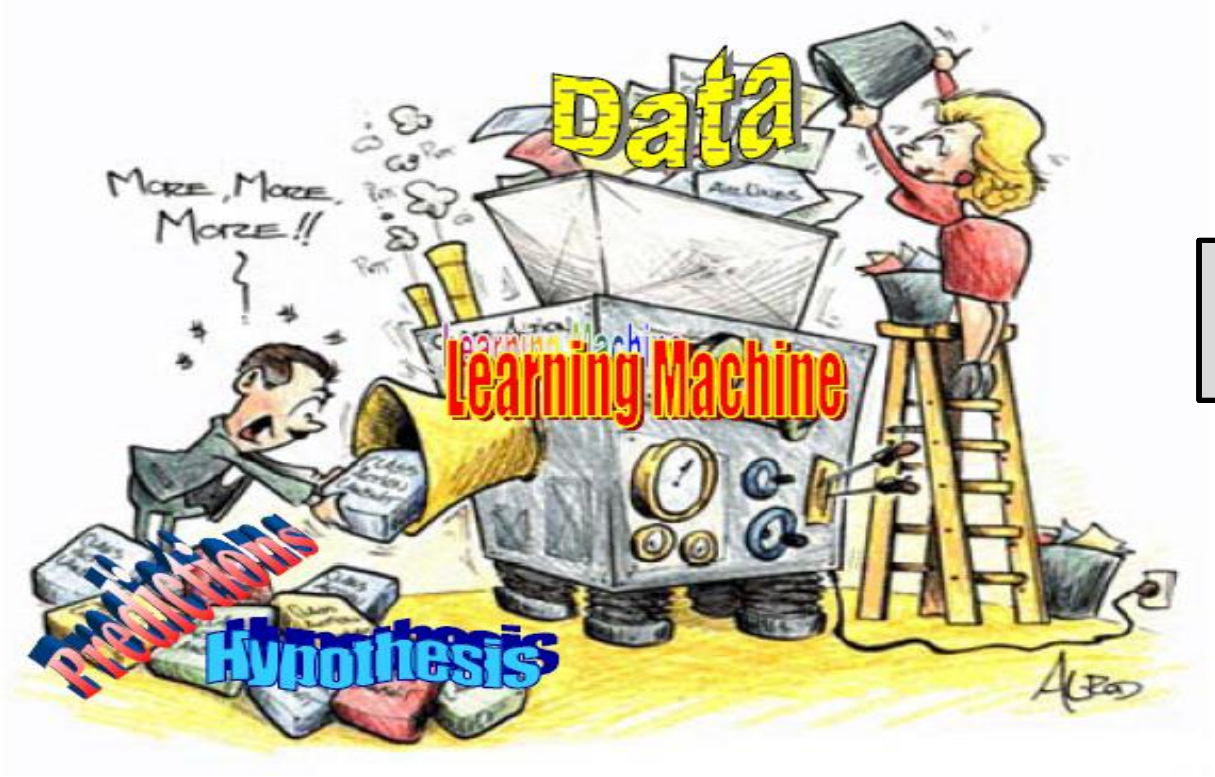

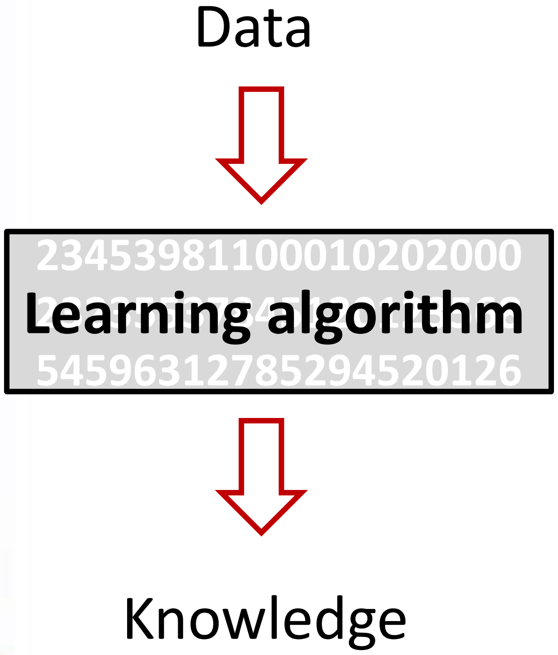

Data

Learning algorithm

Knowledge

What is machine learning?

- Study of algorithms that
  - (automatically) improve their <u>performance</u>
  - at some <u>task</u>
  - with <u>experience</u>

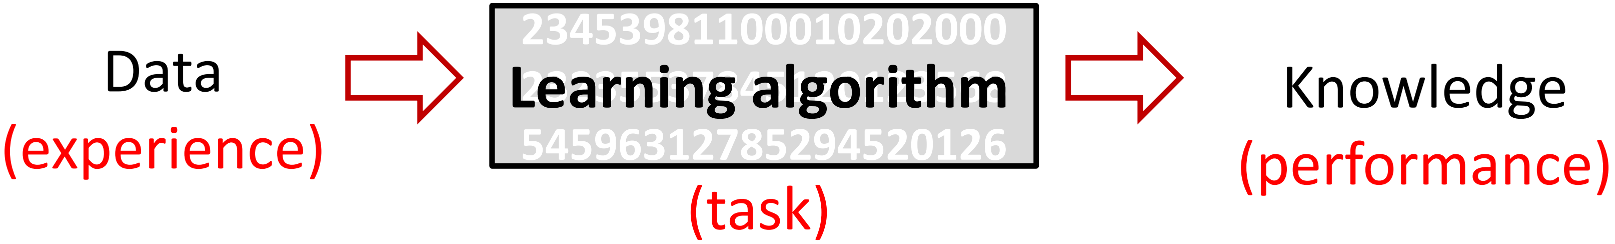

Data
(experience)

Learning algorithm
(task)

Knowledge
(performance)

(Statistical) Machine Learning in AI

pre-computer age

thoughts on symbolic reasoning

1950s & 1960s: symbolic r ni

representation, sea

1960s & 1970s: suc

1970s & 1980s: expert systems

Great growth ac

s 19

pre-Pearl AI

post-Pearl AI

[Judea Pearl, Turing Award 2011]

- For "innovations that enabled remarkable advances in the partnership between humans and machines that is the foundation of Artificial Intelligence (AI)”
- “His work serves as the standard method for handling uncertainty in computer systems, with applications from medical diagnosis, homeland security and genetic counseling to natural language understanding and mapping gene expression data.”
- “Modern applications of AI, such as robotics, self-driving cars, speech recognition, and machine translation deal with uncertainty. Pearl has been instrumental in supplying the rationale and much valuable technology that allow these applications to flourish.”

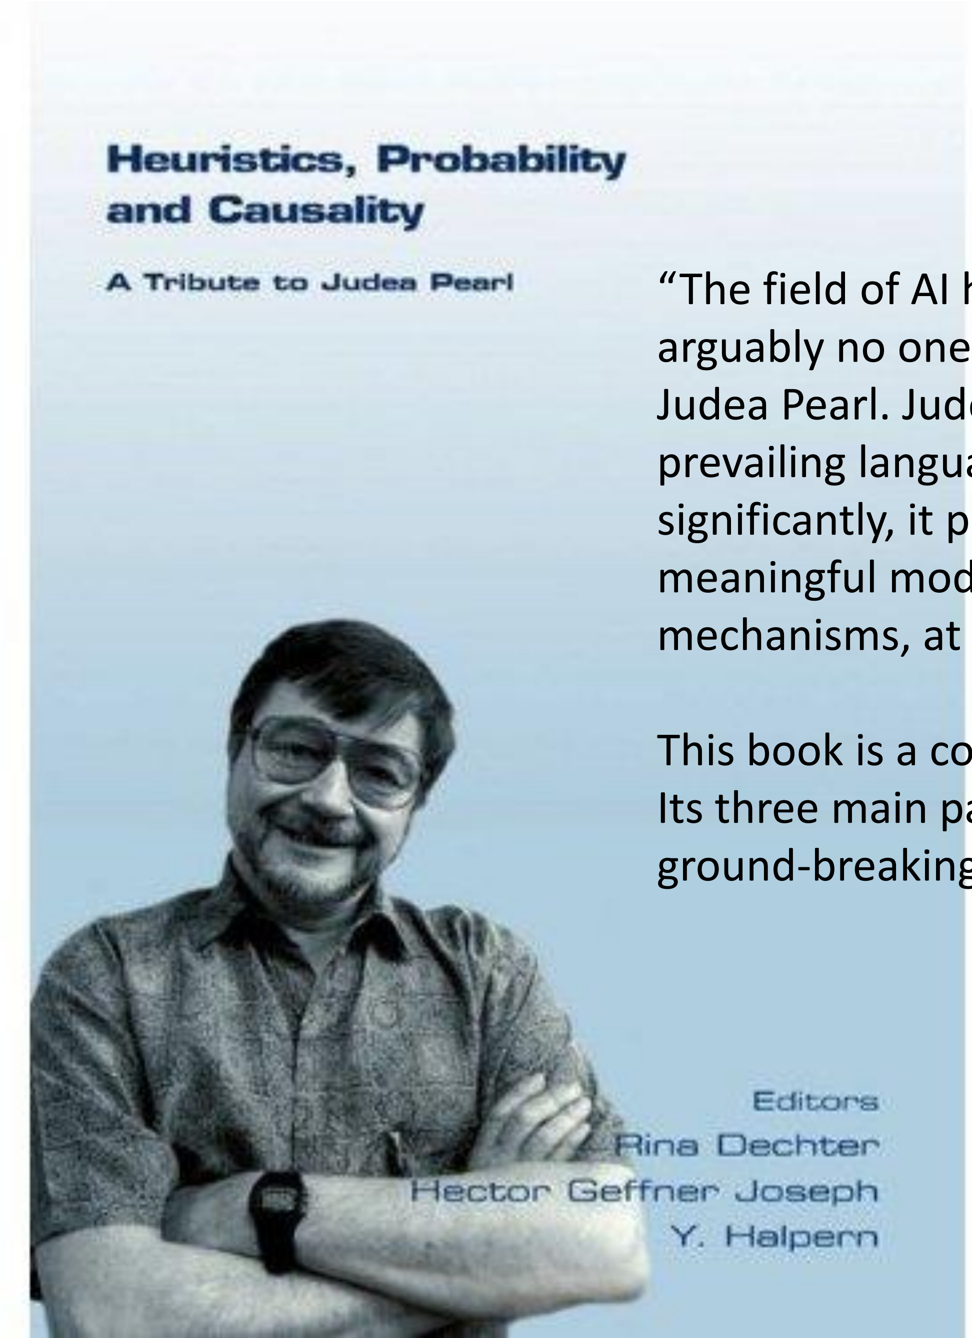

Heuristics, Probability and Causality

A Tribute to Judea Pearl

Editors
Rina Dechter
Hector Geffner Joseph
Y. Halpern

“The field of AI has changed a great deal since the 80s, and arguably no one has played a larger role in that change than Judea Pearl. Judea Pearl's work made probability the prevailing language of modern AI and, perhaps more significantly, it placed the elaboration of crisp and meaningful models, and of effective computational mechanisms, at the center of AI research …”

This book is a collection of articles in honor of Judea Pearl. Its three main parts correspond to the titles of the three ground-breaking books authored by Judea …

Machine learning in Action

- Document classification

Sports
News
Politics
…

Spam Filter

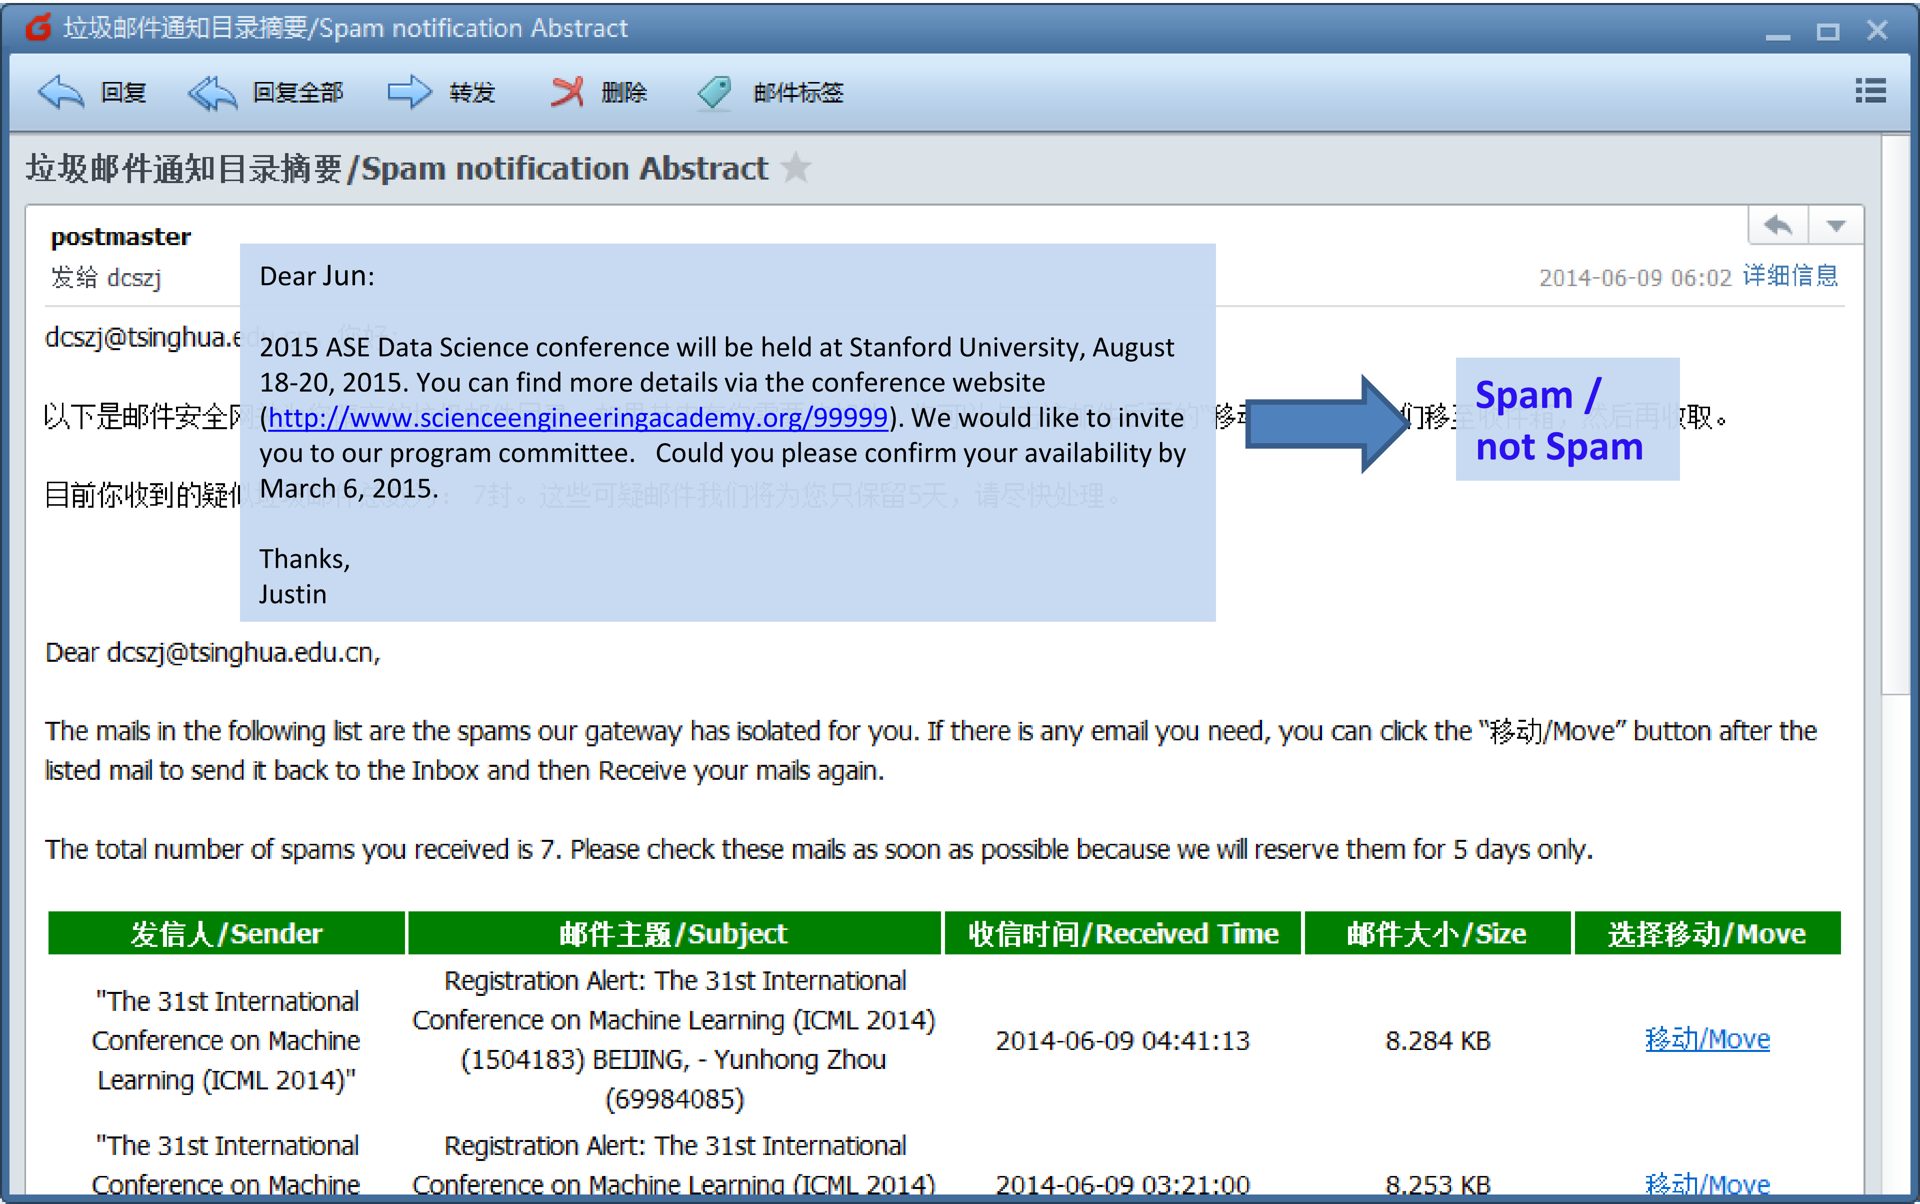

垃圾邮件通知目录摘要/Spam notification Abstract

postmaster

发给 dcszj

2014-06-09 06:02 详细信息

dcszj@tsinghua.e

以下是邮件安全网 … 们移 … 取。

目前你收到的疑似 … 7封。这些可疑邮件我们将为您只保留5天，请尽快处理。

Dear Jun:

2015 ASE Data Science conference will be held at Stanford University, August 18-20, 2015. You can find more details via the conference website (http://www.scienceengineeringacademy.org/99999). We would like to invite you to our program committee.   Could you please confirm your availability by March 6, 2015.

Thanks,
Justin

Spam / not Spam

Dear dcszj@tsinghua.edu.cn,

The mails in the following list are the spams our gateway has isolated for you. If there is any email you need, you can click the "移动/Move" button after the listed mail to send it back to the Inbox and then Receive your mails again.

The total number of spams you received is 7. Please check these mails as soon as possible because we will reserve them for 5 days only.

发信人/Sender	邮件主题/Subject	收信时间/Received Time	邮件大小/Size	选择移动/Move
"The 31st International Conference on Machine Learning (ICML 2014)"	Registration Alert: The 31st International Conference on Machine Learning (ICML 2014) (1504183) BEIJING, - Yunhong Zhou (69984085)	2014-06-09 04:41:13	8.284 KB	移动/Move
"The 31st International Conference on Machine	Registration Alert: The 31st International Conference on Machine Learning (ICML 2014)	2014-06-09 03:21:00	8.253 KB	移动/Move

Regression

- Stock market prediction

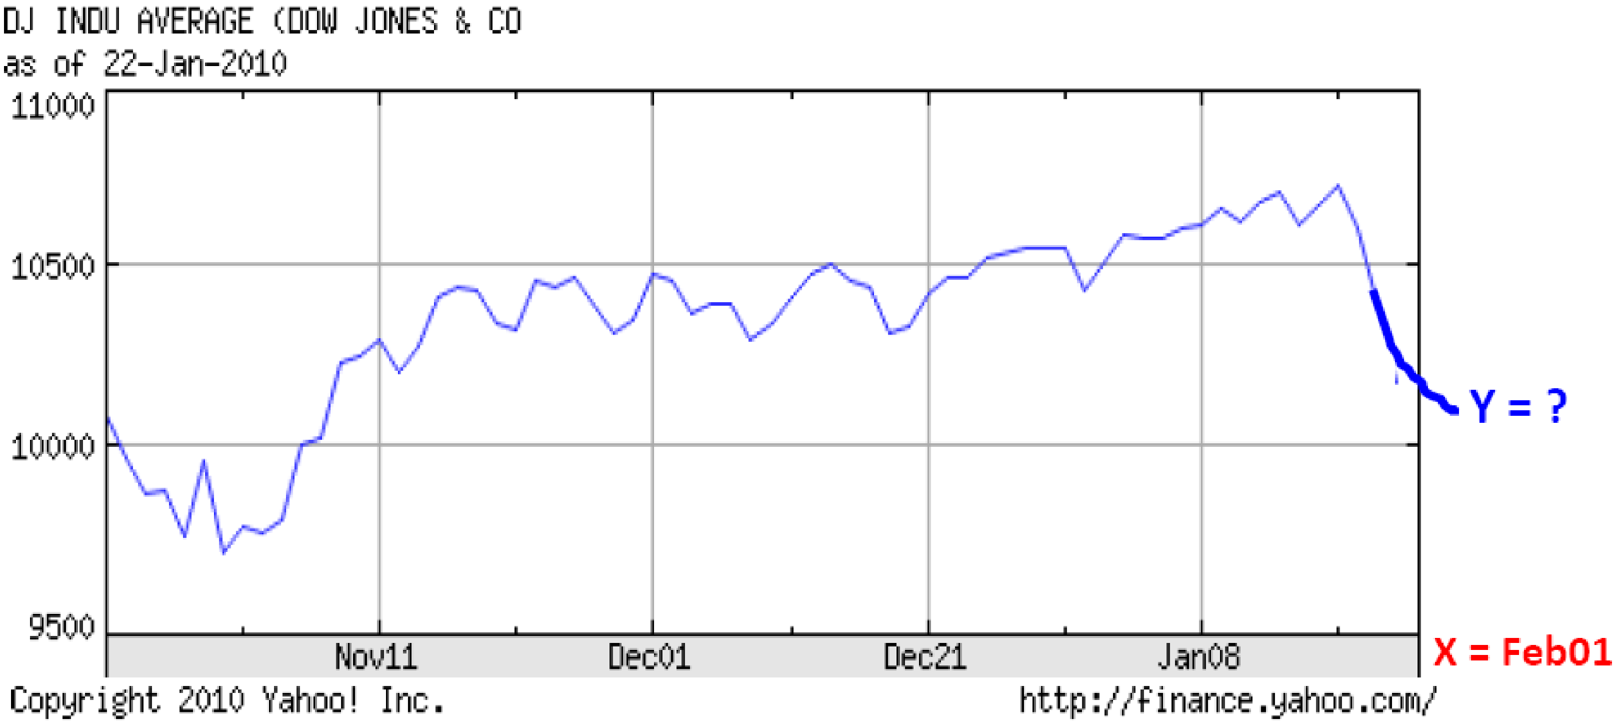

Computer Vision

- Image Classification, Face recognition, Scene understanding, Action/behavior recognition, Image tagging and search, Optical character recognition (OCR)

1000 object classes that we recognize
Images courtesy of ImageNet (http://www.image-net.org/challenges/LSVRC/2010/index)

ImageNet Challenge: 1000 categories, 1.2 million images for training

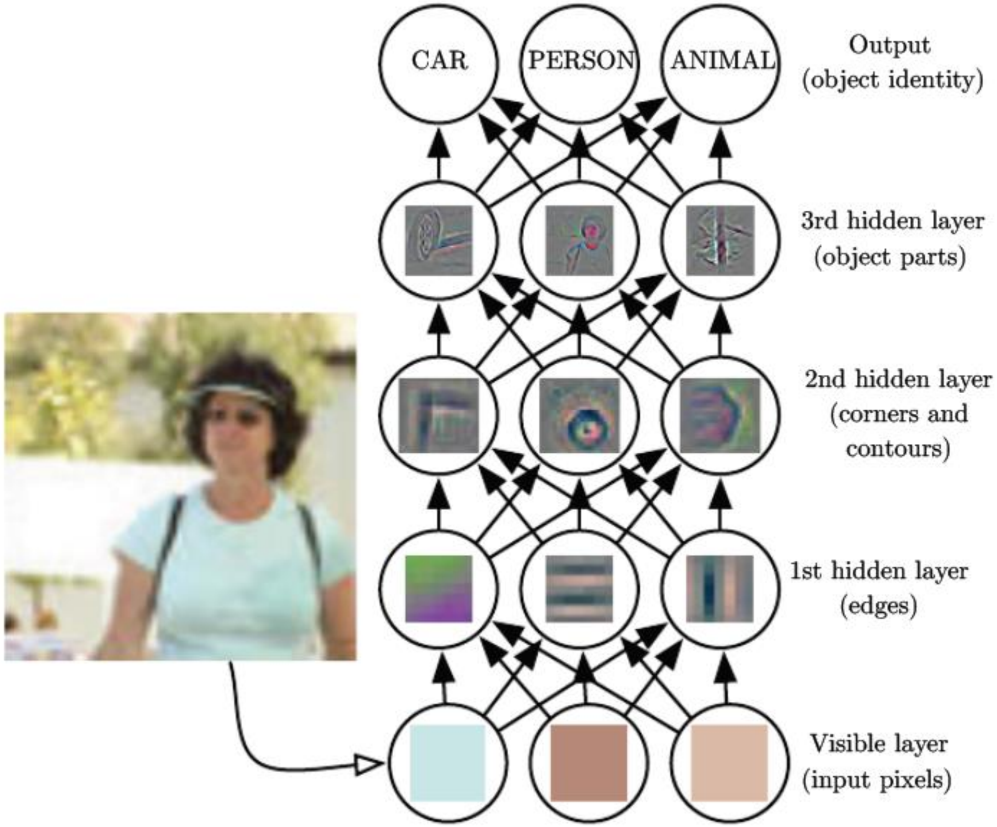
CAR
PERSON
ANIMAL
Output (object identity)
3rd hidden layer (object parts)
2nd hidden layer (corners and contours)
1st hidden layer (edges)
Visible layer (input pixels)

Speech Recognition

- A classic problem in AI, very difficult!
  - “Let’s talk about how to wreck a nice beach”
  - small vocabulary is – easy
  - challenges: large vocabulary, noise, accent, semantics

iPhone 4S
“Find me a French restaurant in Bethesda Maryland”
I found 13 French restaurants... 7 of them are in Bethesda.
Mon Ami Gabi
La Madeleine French Bakery & Cafe

Natural Language Processing

- Machine translation, Information Extraction, Information Retrieval, question answering, Text classification, spam filtering, etc….

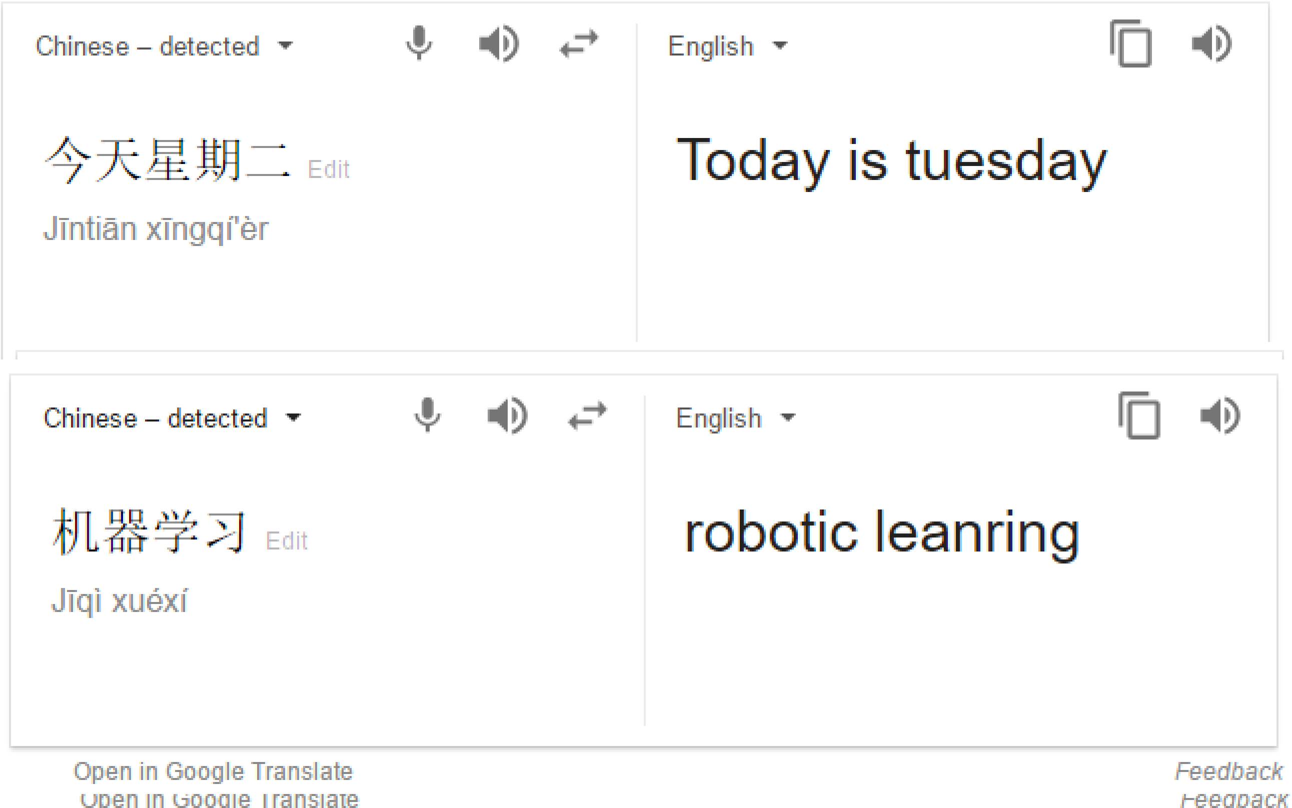

Chinese – detected

今天星期二 Edit

Jīntiān xīngqí'èr

English

Today is tuesday

Chinese – detected

机器学习 Edit

Jīqì xuéxí

English

robotic leanring

Open in Google Translate

Feedback

Open in Google Translate
Feedback

Natural Language Processing

- Machine translation, Information Extraction, Information Retrieval, question answering, Text classification, spam filtering, etc….

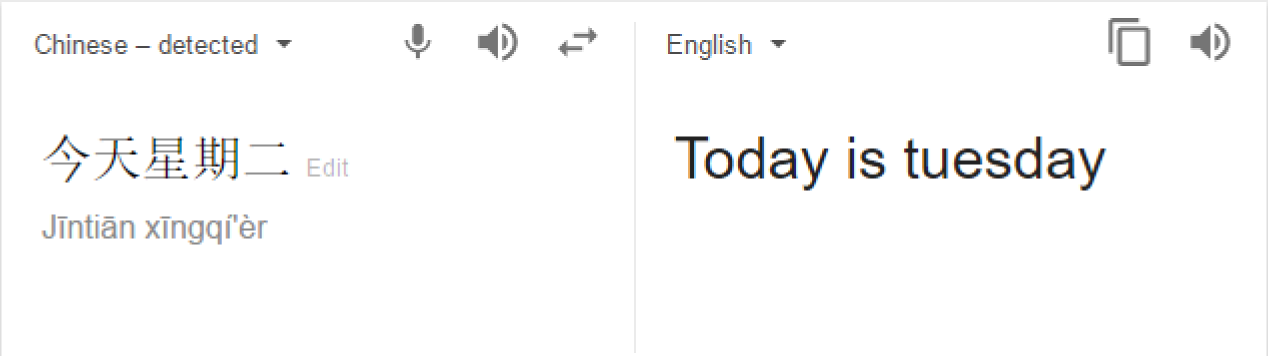

Chinese – detected

English

今天星期二 Edit

Jīntiān xīngqí'èr

Today is tuesday

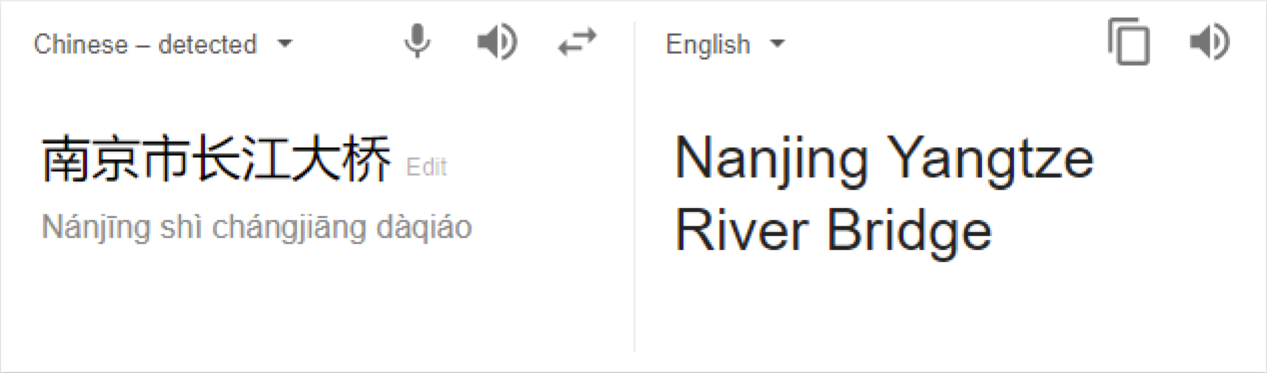

Chinese – detected

English

南京市长江大桥 Edit

Nánjīng shì chángjiāng dàqiáo

Nanjing Yangtze River Bridge

Control

- Cars navigating on their own

  - DAPA urban challenge

  - Tsinghua Mobile Robot V (THMR-V):

- The first license, Nevada, 2012
- Nevada, Florida, California, Michigan, allows testing on public roads

Control (cont'd)

- How the Google car works

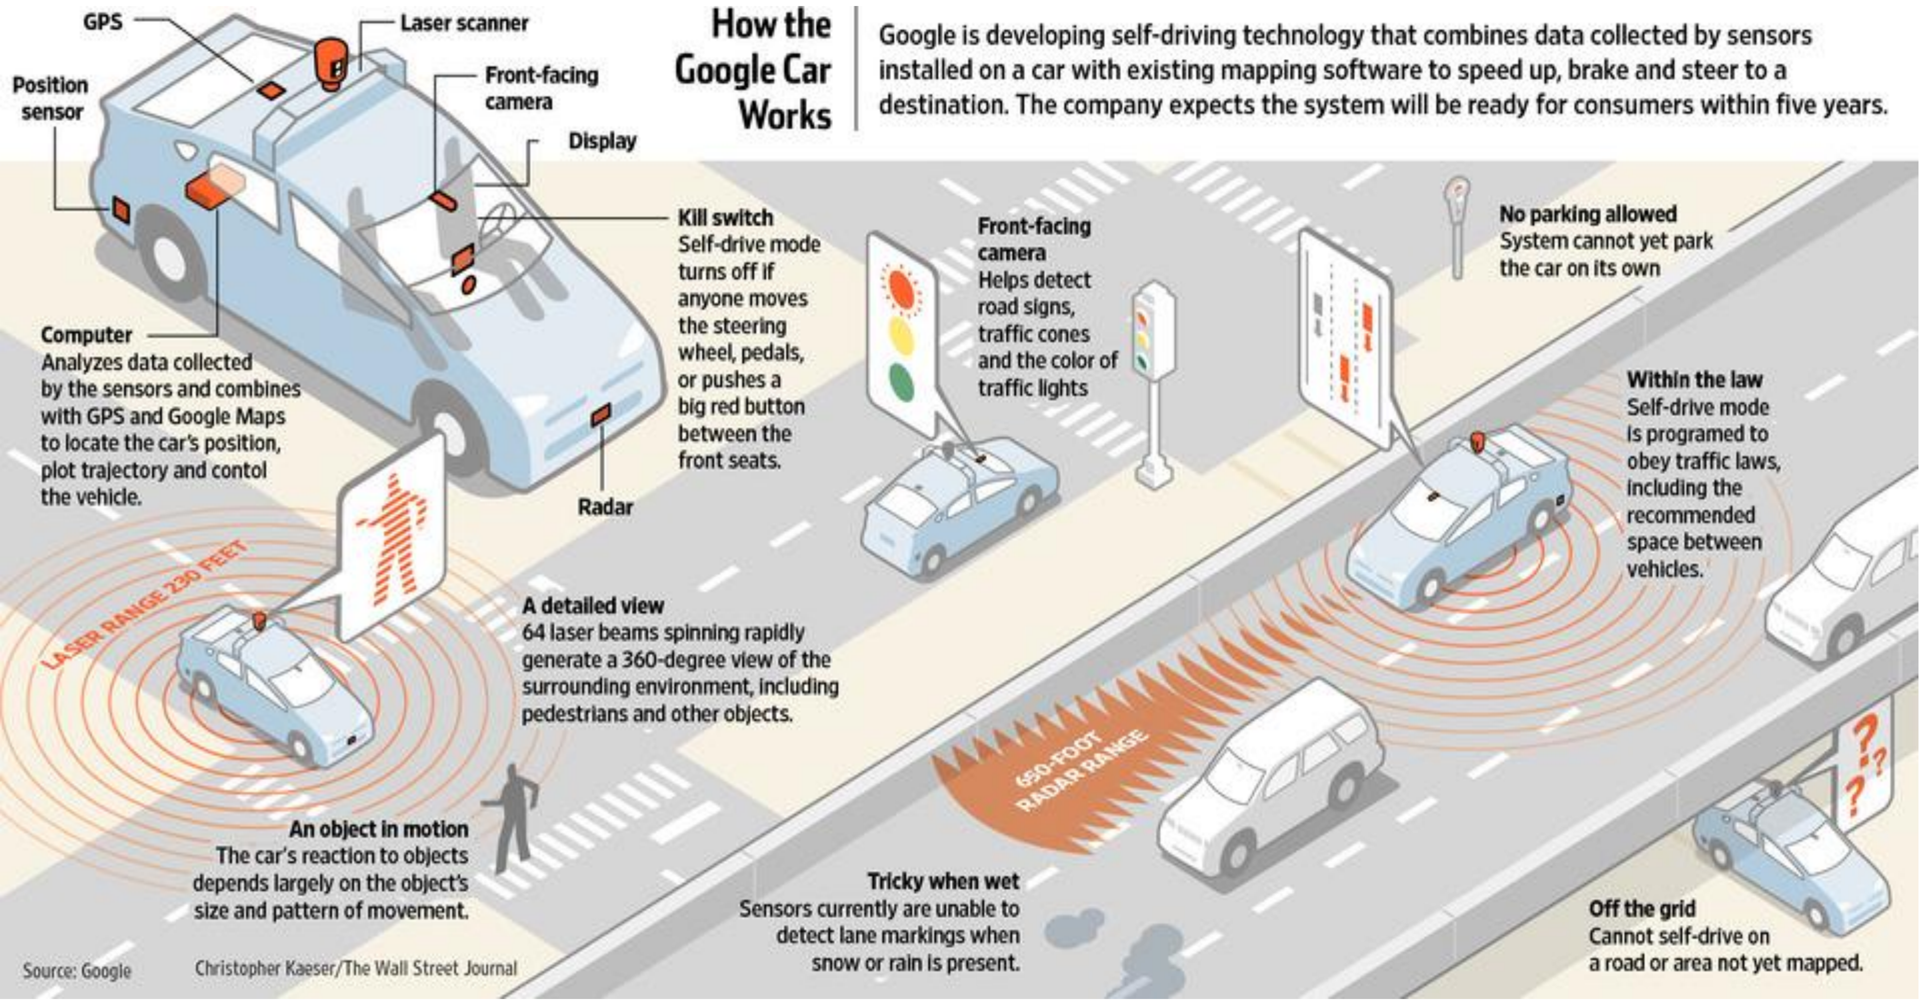

How the Google Car Works

Google is developing self-driving technology that combines data collected by sensors installed on a car with existing mapping software to speed up, brake and steer to a destination. The company expects the system will be ready for consumers within five years.

Computer Analyzes data collected by the sensors and combines with GPS and Google Maps to locate the car's position, plot trajectory and contol the vehicle.

Kill switch Self-drive mode turns off if anyone moves the steering wheel, pedals, or pushes a big red button between the front seats.

Front-facing camera Helps detect road signs, traffic cones and the color of traffic lights

No parking allowed System cannot yet park the car on its own

Within the law Self-drive mode is programed to obey traffic laws, including the recommended space between vehicles.

A detailed view 64 laser beams spinning rapidly generate a 360-degree view of the surrounding environment, including pedestrians and other objects.

An object in motion The car's reaction to objects depends largely on the object's size and pattern of movement.

Tricky when wet Sensors currently are unable to detect lane markings when snow or rain is present.

Off the grid Cannot self-drive on a road or area not yet mapped.

Source: Google

Christopher Kaeser/The Wall Street Journal

AlphaGO

- March, 2016: AlaphaGO beats Sedol Lee at 4:1

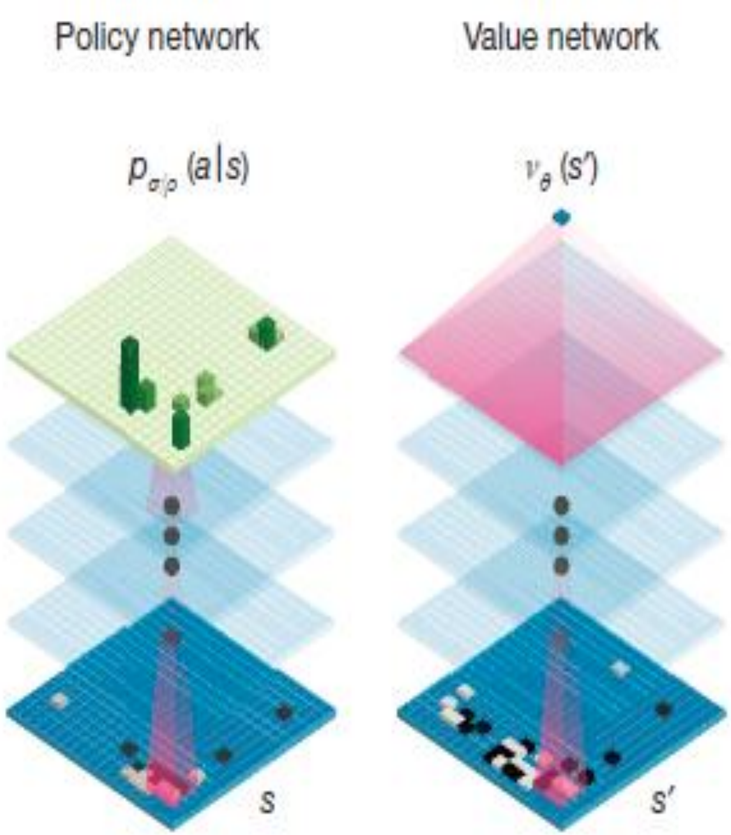

Alpha

- June, 2016: Alpha beats Gene Lee in combat simulation

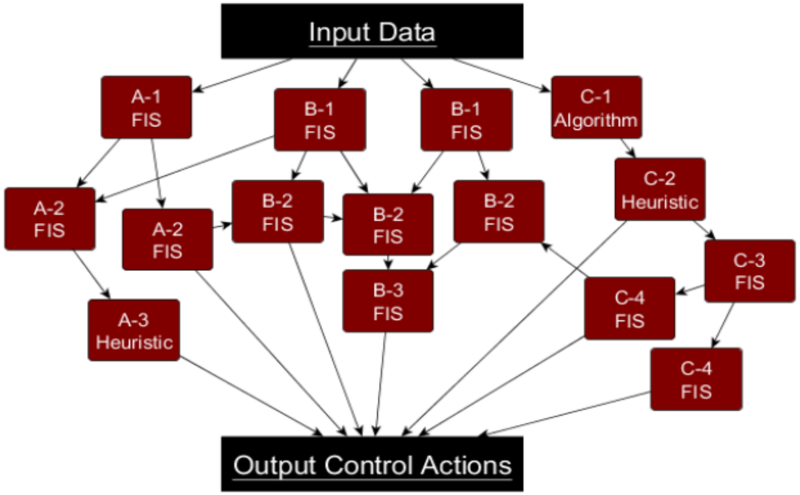

Input Data

A-1 FIS

B-1 FIS

B-1 FIS

C-1 Algorithm

A-2 FIS

A-2 FIS

B-2 FIS

B-2 FIS

B-2 FIS

C-2 Heuristic

C-3 FIS

B-3 FIS

C-4 FIS

A-3 Heuristic

C-4 FIS

Output Control Actions

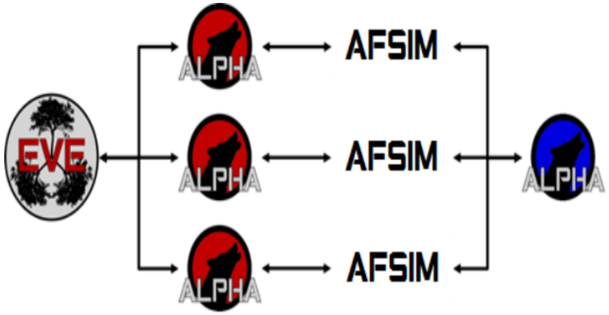

Science

- Decoding thoughts from brain activity

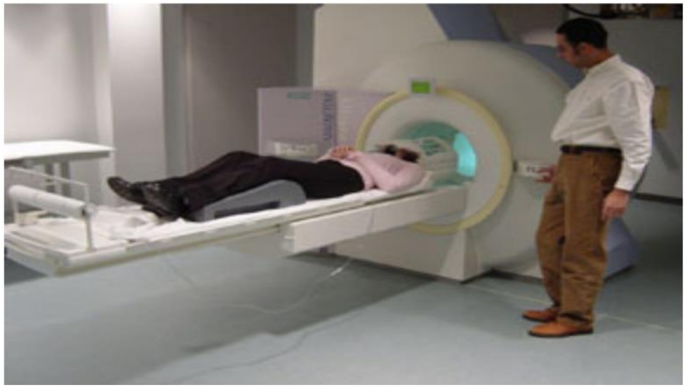

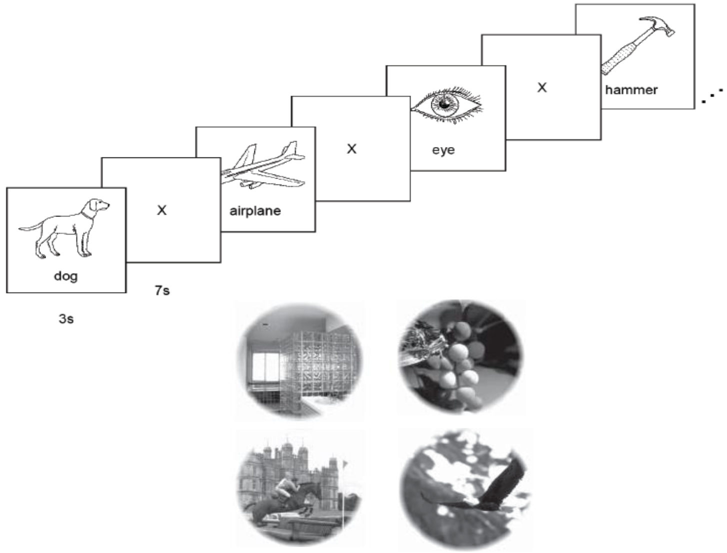

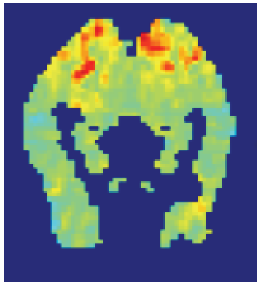

Tool
Animal
…

[Mitchell et al, Science 2008]
[Kay et al., Nature, 2008]

Science (cont’d)

- Bayesian models of inductive learning and reasoning [Tenenbaum et al., Science 2011]
  - Challenge:
    - *How can people generalize well from sparse, noisy, and ambiguous data?*
  - Hypothesis:
    - *If the mind goes beyond the data given, some more abstract background knowledge must generate and delimit the possible hypotheses*
  - Bayesian models make structured abstract knowledge and statistical inference cooperate
  - Examples
    - Word learning [Xu & Tenenbaum, Psychol. Rev. 2007]
    - Causal relation learning [Griffiths & Tenenbaum, 2005]
    - Human feature learning [Austerweil & Griffiths, NIPS 2009]
    - …

J. Tenenbaum et al., “How to grow a mind: Statistics, Structure, and Abstraction”. Science 331, 1279 (2011)

More others …

- Many more
  - Natural language processing
  - Speech recognition
  - Computer vision
  - Robotics
  - Computational biology
  - Social network analysis
  - Sensor networks
  - Health care
  - Protest ??
  - …

Machine learning in Action

- Machine learning for protest?

CMU ML students and post-docs at G-20 Pittsburgh Summit 2009

Machine Learning – practice

document classification

decoding brain signal

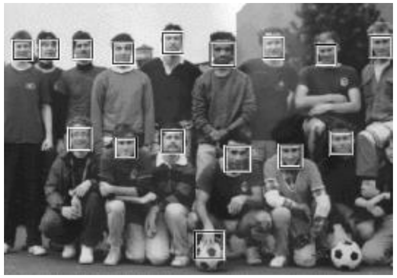

face recognition

robot control

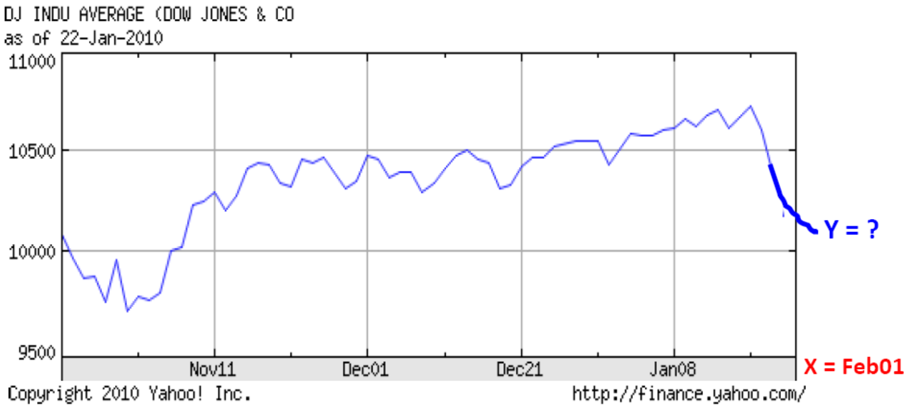

stock market prediction

Machine Learning – theory

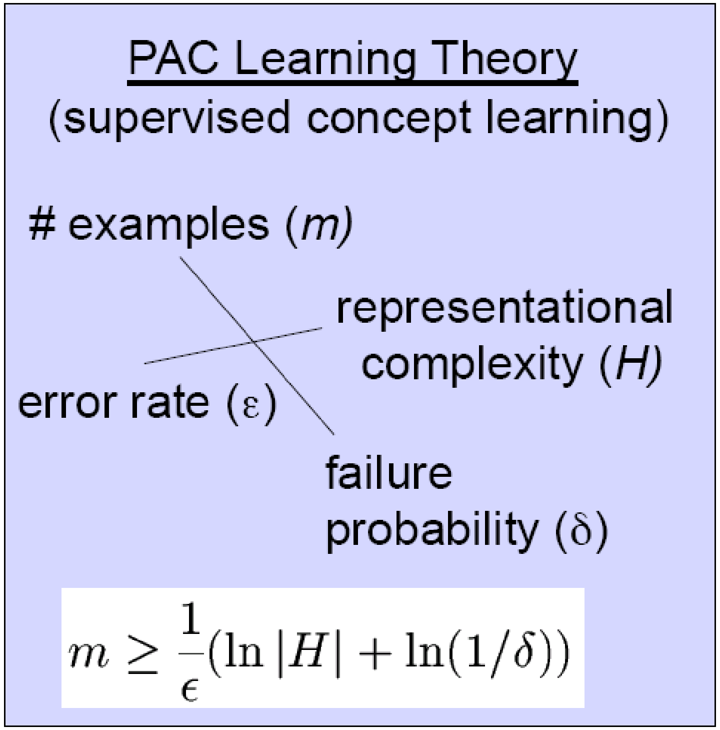

PAC Learning Theory

(supervised concept learning)

$$m \geq \frac{1}{\epsilon}(\ln|H| + \ln(1/\delta))$$

Other theories for

- semi-supervised learning
- reinforcement skill learning
- active learning
- …

… also relating to

- # mistakes during training
- asymptotic performance
- convergence rate
- bias, variance tradeoff
- …

[Leslie G. Valiant, 1984; Turing Award, 2010]

“For transformative contributions to the theory of computation, including the theory of probably approximately correct (PAC) learning, the complexity of enumeration and of algebraic computation, and the theory of parallel and distributed computing.”

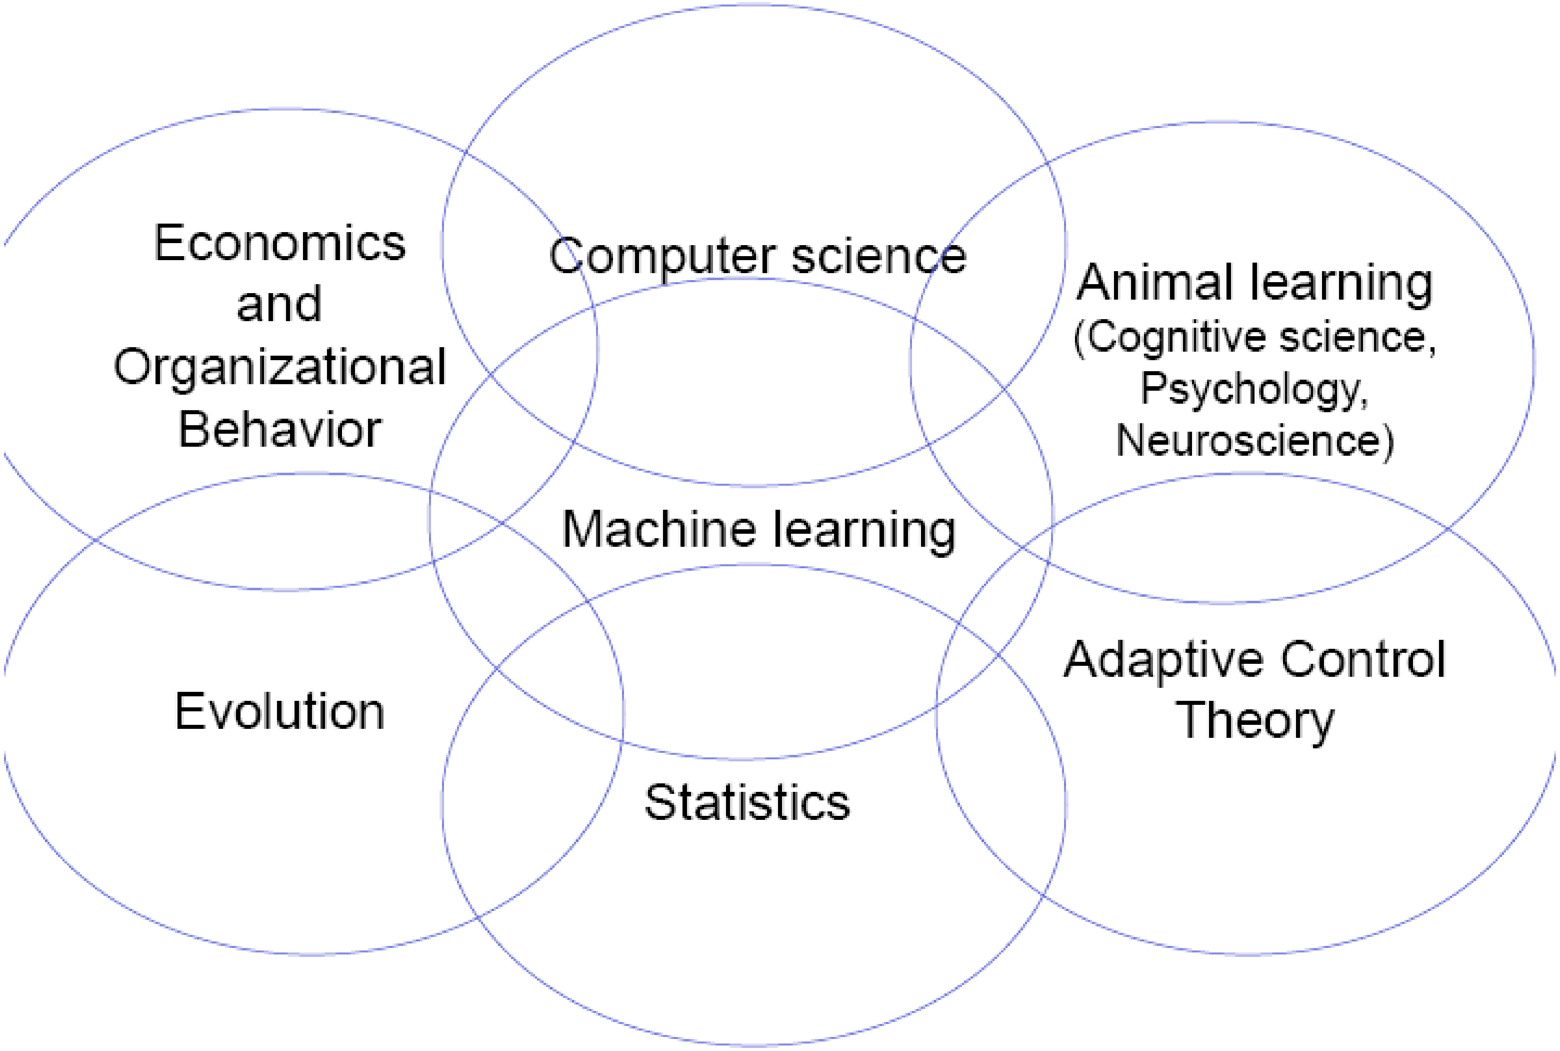

Economics and Organizational Behavior

Computer science

Animal learning (Cognitive science, Psychology, Neuroscience)

Machine learning

Evolution

Adaptive Control Theory

Statistics

Growth of Machine Learning in CS

- Machine learning already the preferred approach to
  - Speech recognition, natural language process
  - Computer vision
  - Medical outcomes analysis
  - Robot control
  - …

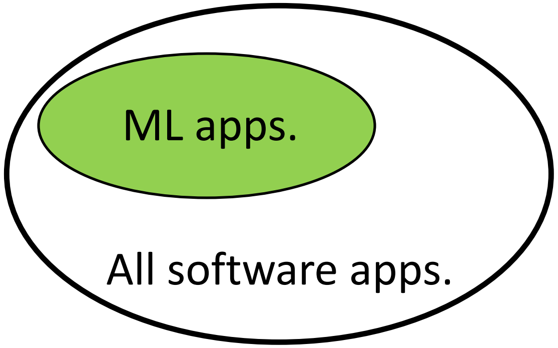

- This ML niche is growing (why?)

Growth of Machine Learning in CS

- Machine learning already the preferred approach to
  - Speech recognition, natu
  - Computer vision
  - Medical outcomes analy
  - Robot control
  - …
- This ML niche is growing
  - Improved machine learning algorithms
  - Increased data capture, networking, new sensors
  - Software too complex to write by hand
  - Demand for self-customization to user, environment

Huge amount of data …

- Web: estimated Google index 45 billion pages
- Transaction data: 5-50 TB/day
- Satellite image feeds: ~1TB/day/satellite
- Biological data: 1-10TB/day/sequencer
- TV: 2TB/day/channel;
- YouTube 4TB/day uploaded
- Photos: 1.5 billion photos/week uploaded

ML has a long way to go …

- Very large-scale learning in rich media

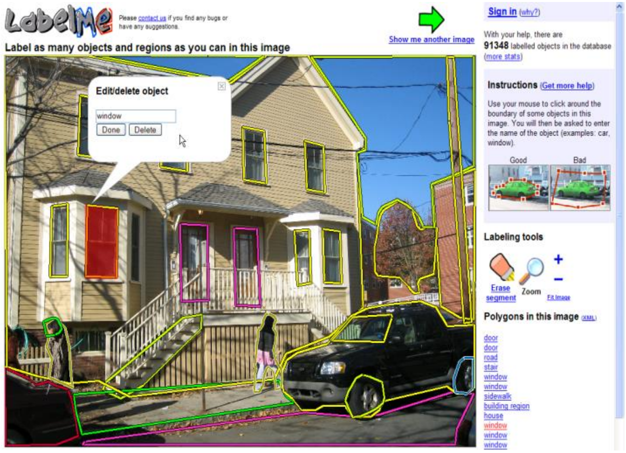

Please contact us if you find any bugs or have any suggestions.

Label as many objects and regions as you can in this image

With your help, there are 91348 labelled objects in the database (more stats)

Instructions (Get more help)

Use your mouse to click around the boundary of some objects in this image. You will then be asked to enter the name of the object (examples: car, window).

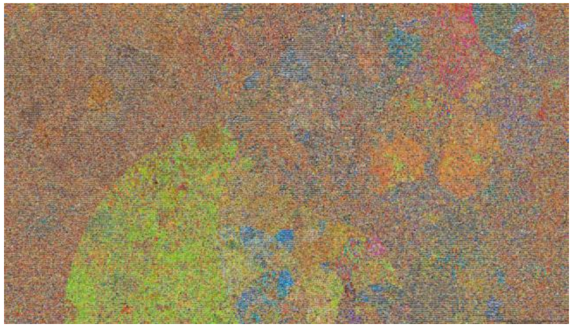

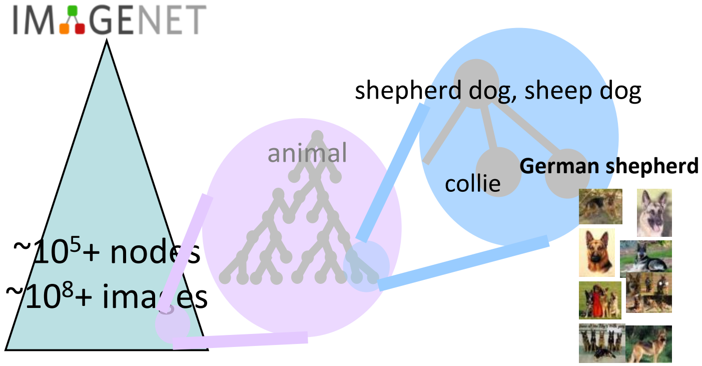

10^5 images 10^{1-2} categories → 10^5 images 10^{2-3} categories → 10^{6-7} images 10^{3-4} categories

Machine Learning Tasks

- Broad categories
  - Supervised learning
    - Classification, Regression
  - Unsupervised learning
    - Density estimation, Clustering, Dimensionality reduction
  - Reinforcement learning
  - Semi-supervised learning
  - Active learning
  - Transfer learning
  - Many more …

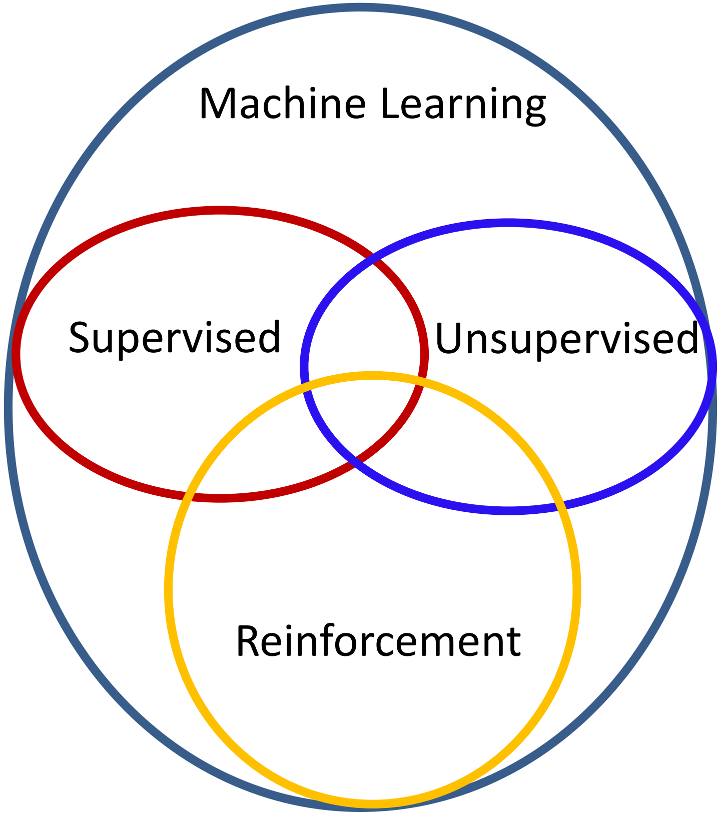

Supervised Learning

- Task: learn a predictive function $h:\ \mathcal{X} \rightarrow \mathcal{Y}$

Feature space $\mathcal{X}$

Label space $\mathcal{Y}$

Words in documents

“Sports”
“News”
“Politics”
…

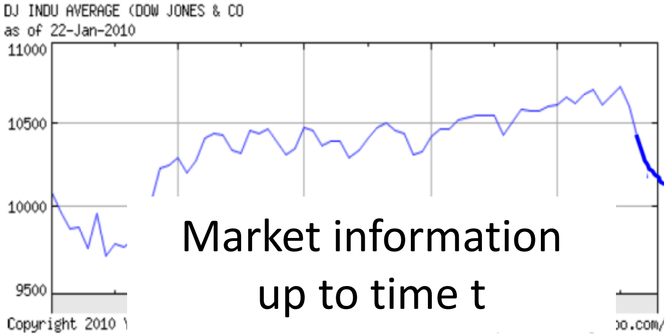

Market information
up to time t

Share price
“$ 20.50”

- “Experience” or training data:

$$\{< x_d, y_d >\}_{d=1}^{D},\ \ x_d \in \mathcal{X}, y_d \in \mathcal{Y}$$

Supervised Learning – classification

Feature space $\mathcal{X}$

Label space $\mathcal{Y}$

Words in documents

⇨

“Sports”
“News”
“Politics”
…

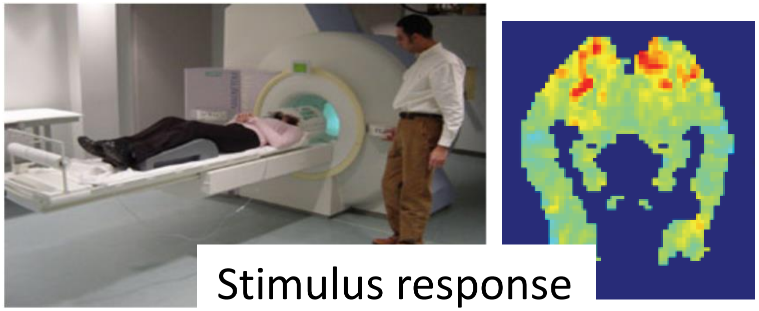

Stimulus response

⇨

“Tool”
“Animal”
…

Discrete Labels

Supervised Learning – regression

Feature space $\mathcal{X}$

Label space $\mathcal{Y}$

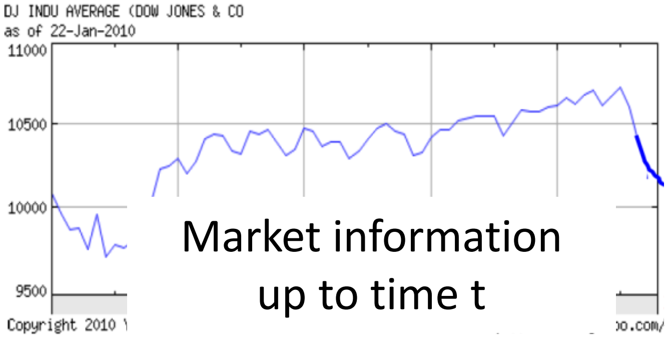

Market information up to time t

Share price "$ 20.50"

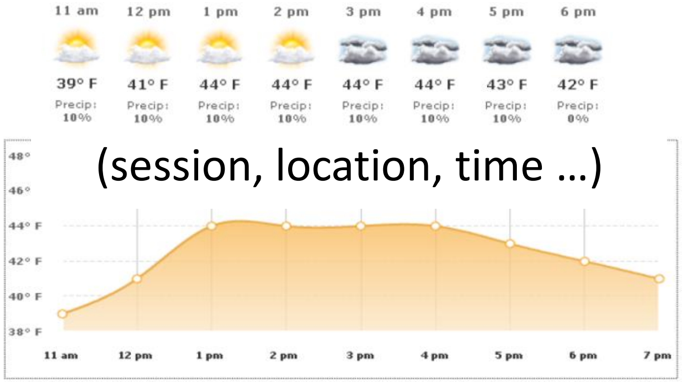

(session, location, time …)

Temperature "42° F"

Continuous Labels

How to learn a classifier?

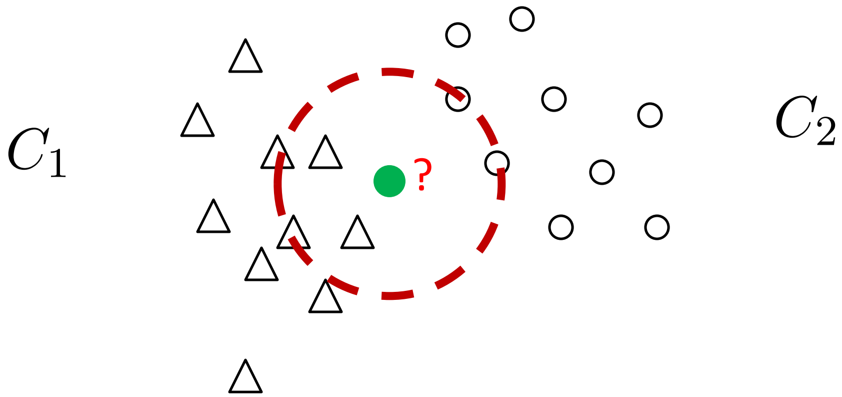

K-NN: a Non-parametric approach

Distance metric matters!

How to learn a classifier?

Parametric (model-based) approaches:

$$g(x) = 0,\ where\ g(x) = w^\top x + w_0$$

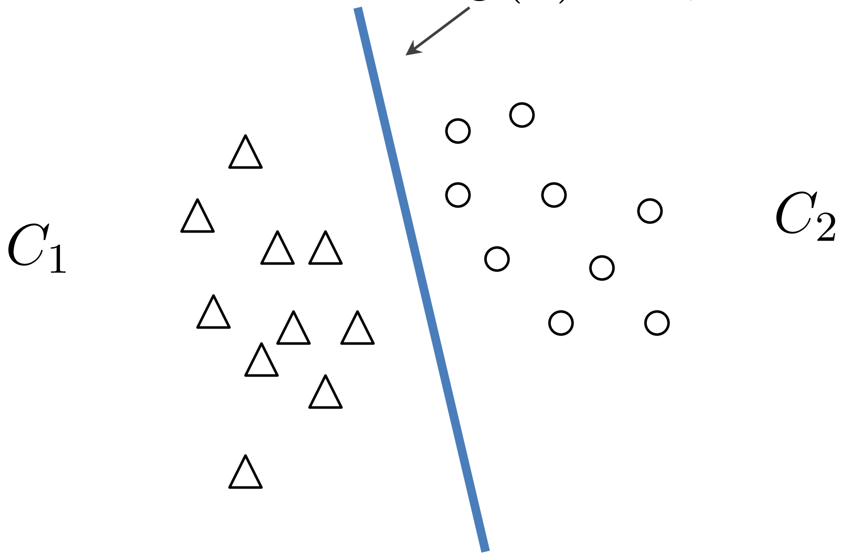

a good decision boundary

$$y^* = \begin{cases} C_1 & if\ \ g(x) > 0 \\ C_2 & if\ \ g(x) < 0 \end{cases}$$

How to learn a classifier?

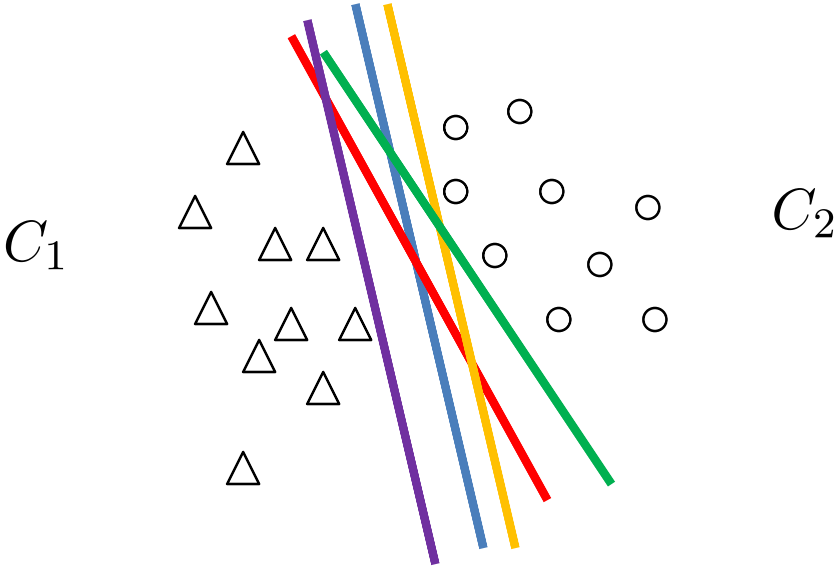

Many good decision boundaries

which one should we choose?

How to learn a classifier?

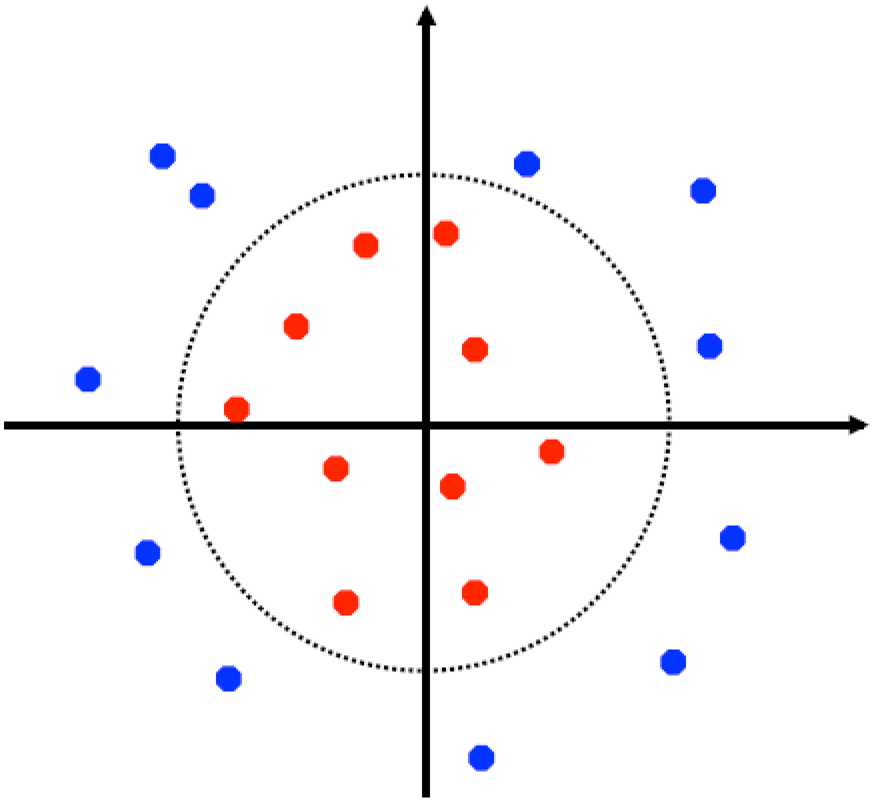

How about non-linearity?

How to learn a classifier?

- 2D mapping is insufficient

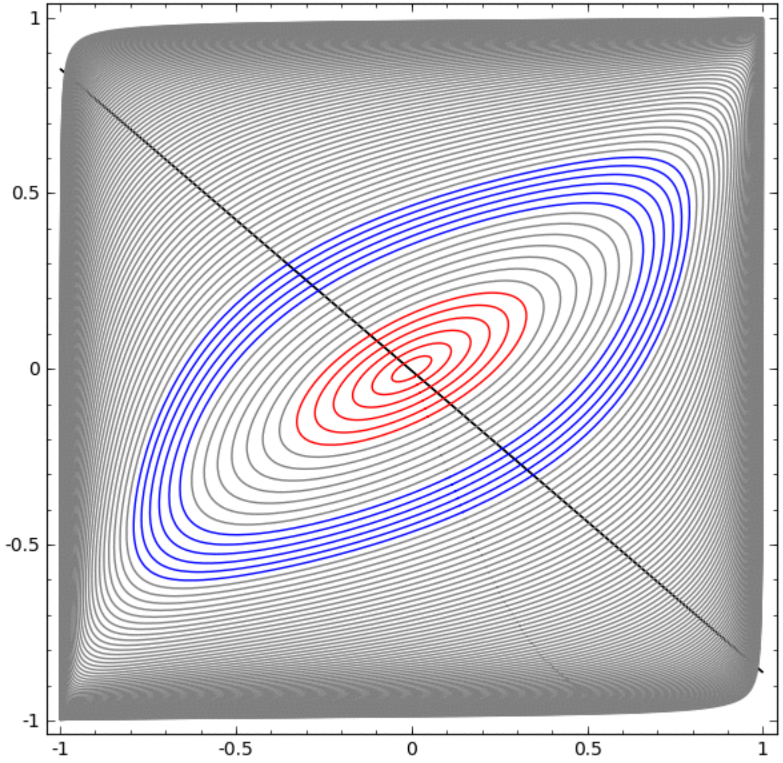

How to learn a classifier?

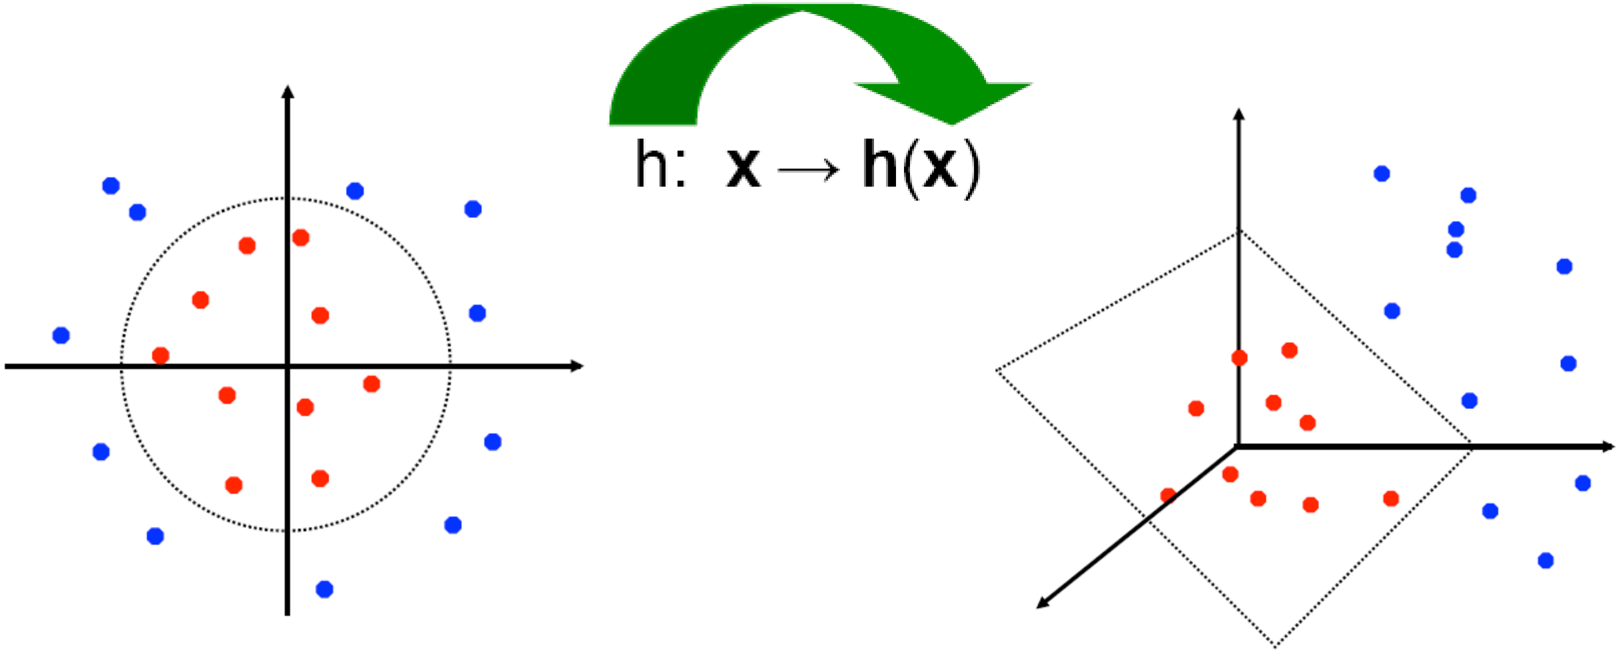

How about non-linearity?

The higher dimension, the better?

How to learn a classifier?

- Curse of dimensionality
  - A high dimensional space is always almost empty

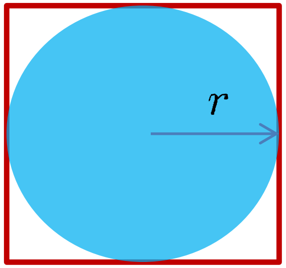

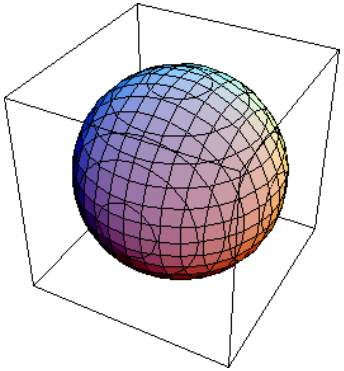

d dimensional space

$$\frac{\pi r^2}{(2r)^2} = \frac{\pi}{4}$$

$$\frac{\frac{2r^3\pi^{3/2}}{3\Gamma(3/2)}}{(2r)^3} = \frac{\pi^{3/2}}{12\Gamma(3/2)}$$

$$\frac{\frac{2r^d\pi^{d/2}}{d\Gamma(d/2)}}{(2r)^d} = \overset{\overset{0}{\uparrow\; d \to \infty}}{\frac{\pi^{d/2}}{d2^{d-1}\Gamma(d/2)}}$$

How to learn a classifier?

- Curse of dimensionality
  - A high dimensional space is always almost empty

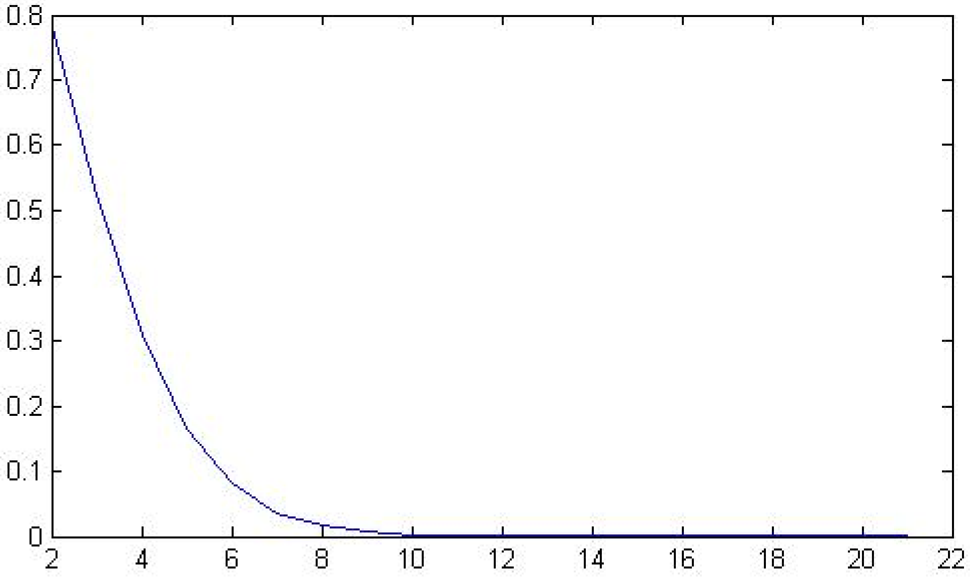

when one wants to learn pattern from data in high dimensions no matter how much data you have it always seems less!

How to learn a classifier?

- Curse of dimensionality

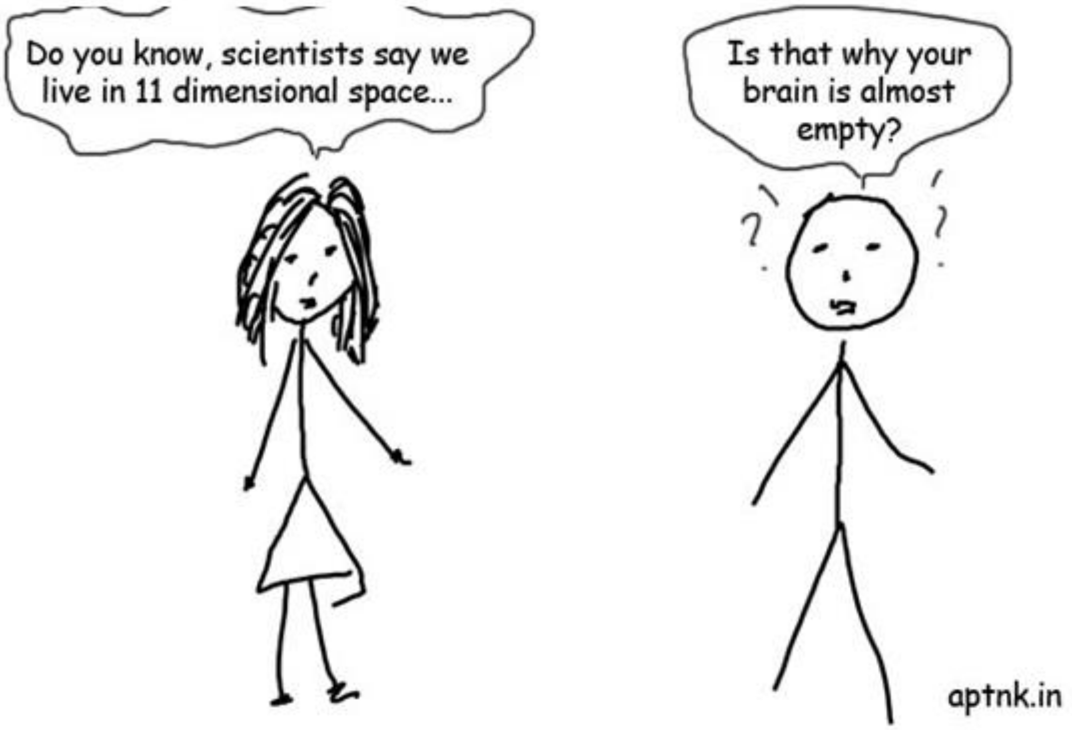
Do you know, scientists say we live in 11 dimensional space...
Is that why your brain is almost empty?
? ?
aptnk.in

when one wants to learn pattern from data in high dimensions no matter how much data you have it always seems less! A high dimensional space is always almost empty

How to learn a classifier?

- Curse of dimensionality
  - A high dimensional space is always almost empty
  - … in high dimensions no matter how much data you have it always seems less!

- The blessing of dimensionality
  - *… real data highly concentrate on low-dimensional, sparse, or degenerate structures in the high-dimensional space.*

- But no free lunch: *Gross errors and irrelevant measurements are now ubiquitous in massive cheap data.*

How to learn a classifier?

- The blessing of dimensionality
  - *… real data highly concentrate on low-dimensional, sparse, or degenerate structures in the high-dimensional space.*

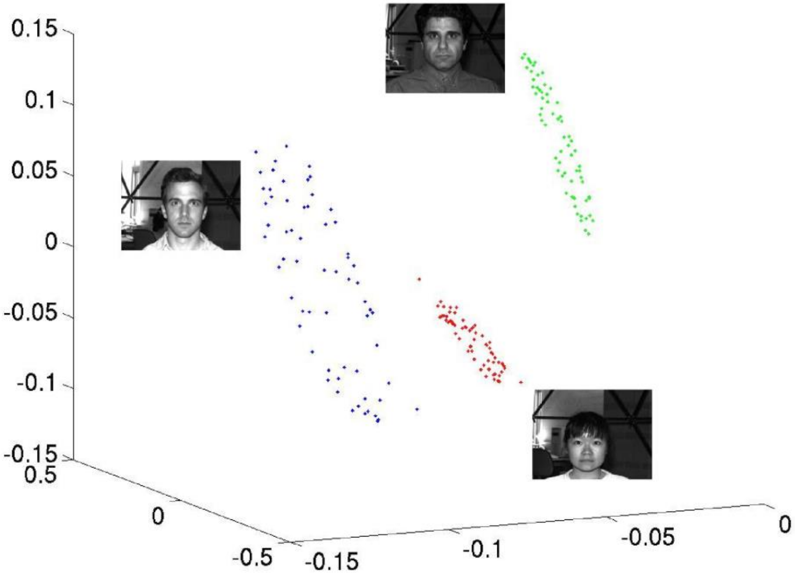

Images of the same face under varying illumination lie approximately on a low (nine)-dimensional subspace, known as the harmonic plane [Basri & Jacobs, PAMI, 2003].

How to learn a classifier?

- Support vector machines (SVM) – basics
  - SVM is among the most popular/successful classifiers
  - It provides a *principled way* to learn a *robust* classifier (i.e., a *decision boundary*)

- SVM
  - chooses the one with *maximum margin principle*
  - has sound *theoretical guarantee*
  - extends to *nonlinear decision boundary* by using *kernel* trick
  - learning problem efficiently solved using convex optimization techniques

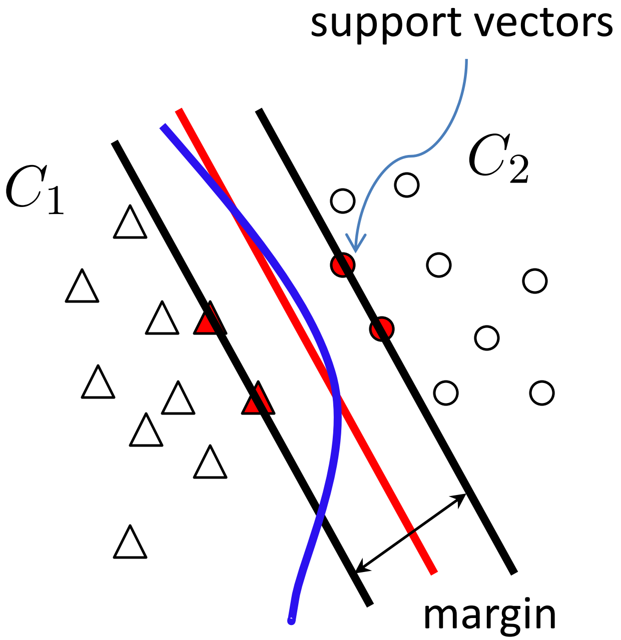

How to learn a classifier?

- Support vector machines (SVM) – demo

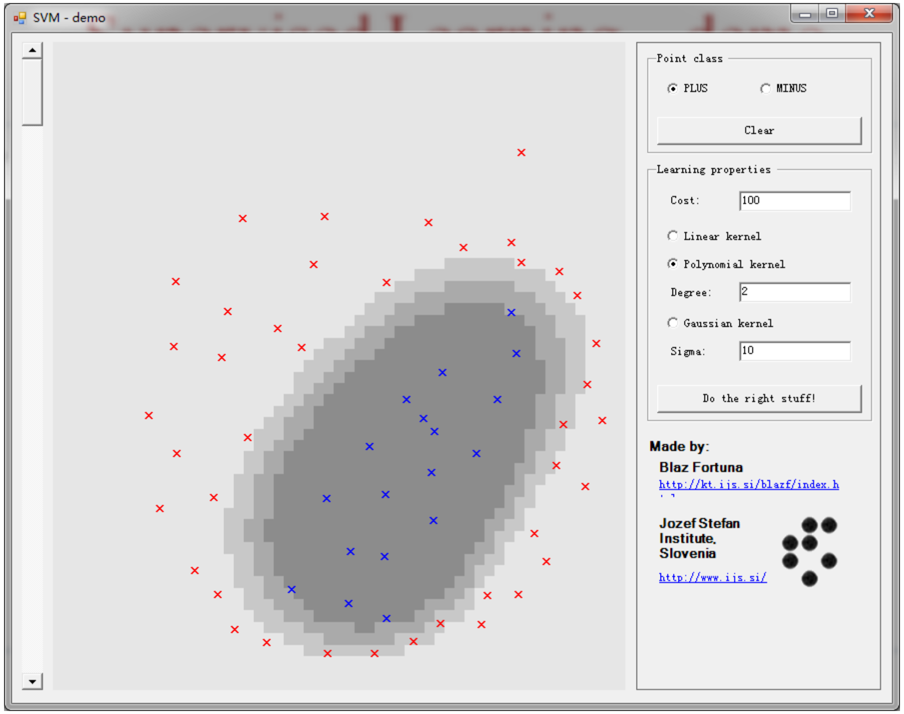

Point class

● PLUS ○ MINUS

Learning properties

Cost: 100

○ Linear kernel

● Polynomial kernel

Degree: 2

○ Gaussian kernel

Sigma: 10

Made by:

Blaz Fortuna
http://kt.ijs.si/blazf/index.h

Jozef Stefan Institute, Slovenia
http://www.ijs.si/

Good ToolKits: [1] SVM-Light: http://svmlight.joachims.org/
[2] LibSVM: http://www.csie.ntu.edu.tw/~cjlin/libsvm/

How to learn a classifier?

- Naïve Bayes classifier – basics
  - an representative method from the very important family of *probabilistic graphical models* and *Bayesian methods*

A joint distribution： $p(x, y) = p(y)p(x|y)$

Inference using Bayes rule：

prior

likelihood

$$p(y|x) = \frac{p(x, y)}{p(x)} = \frac{p(y)p(x|y)}{p(x)}$$

evidence

Prediction rule： $y^* = \arg\max_{y \in \mathcal{Y}} p(y|x)$

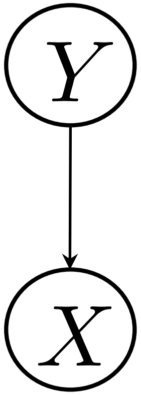

  - fundamental building blocks for *Bayesian networks*
  - nice illustrative example of Bayesian methods

How to learn a classifier?

- Naïve Bayes classifier – basics
  - binary example

$$g(x) \triangleq \log \frac{p(Y=C_1|x)}{p(Y=C_2|x)} = 0$$

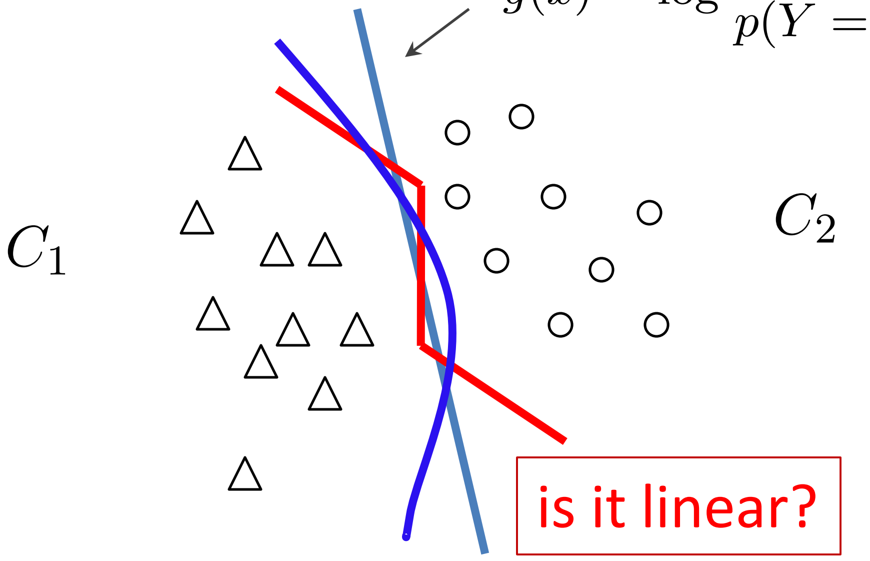

is it linear?

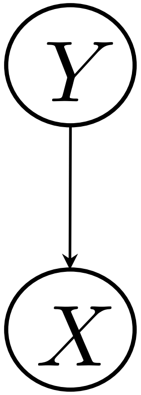

$$y^* = \begin{cases} C_1 & if\ p(Y=C_1|x) > 0.5 \\ C_2 & if\ p(Y=C_1|x) < 0.5 \end{cases}$$

It is for generalized linear models (GLMs)

How to learn a classifier?

- Many other classifiers
  - K-nearest neighbors
  - Decision trees
  - Logistic regression
  - Boosting
  - Random forests
  - Mixture of experts
  - Maximum entropy discrimination (a nice combination of max-margin learning and Bayesian methods)
  - …

Advice #1:

All models are wrong, but some are useful. – G.E.P. Box

Are complicated models preferred?

- A simple curve fitting task

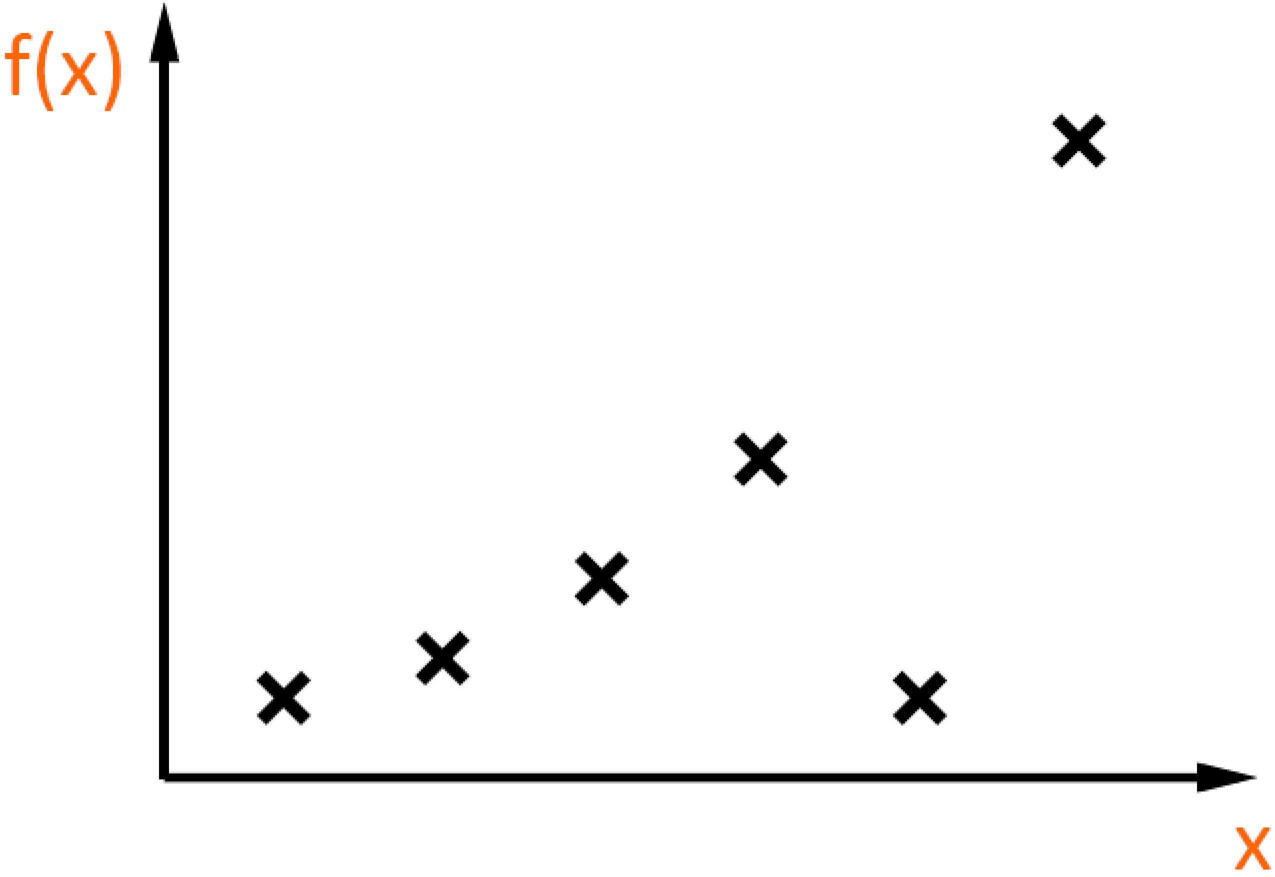

Are complicated models preferred?

- Order = 1

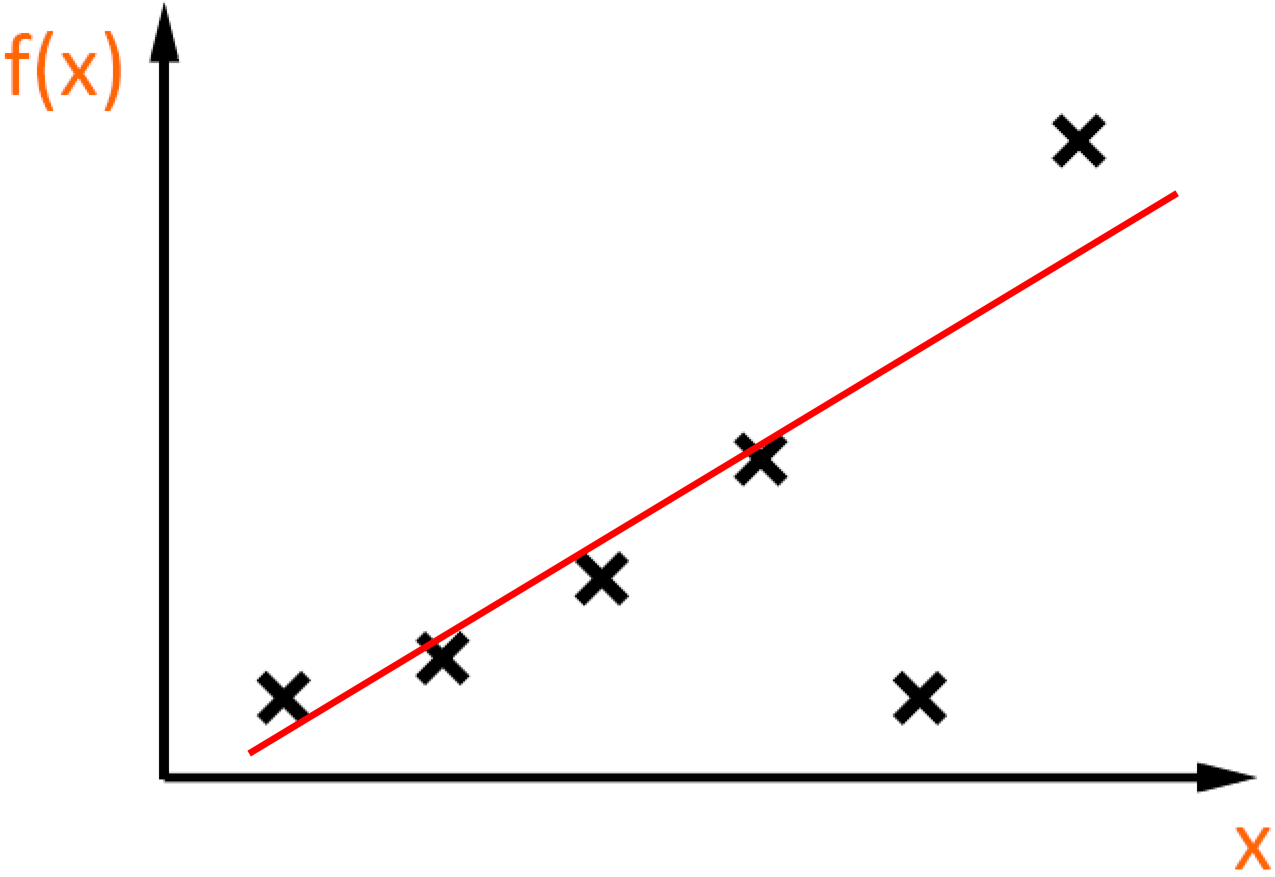

Are complicated models preferred?

- Order = 2

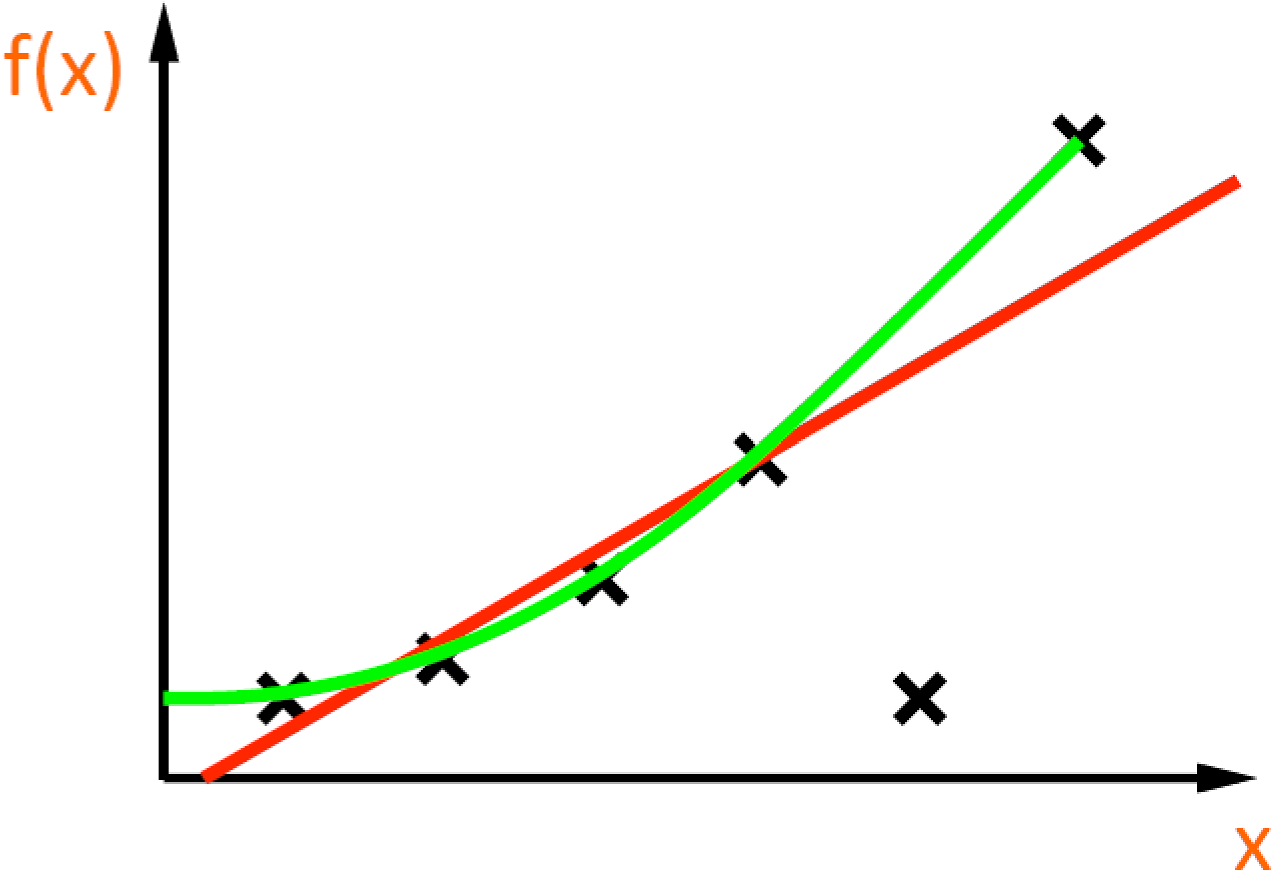

Are complicated models preferred?

- Order = 3

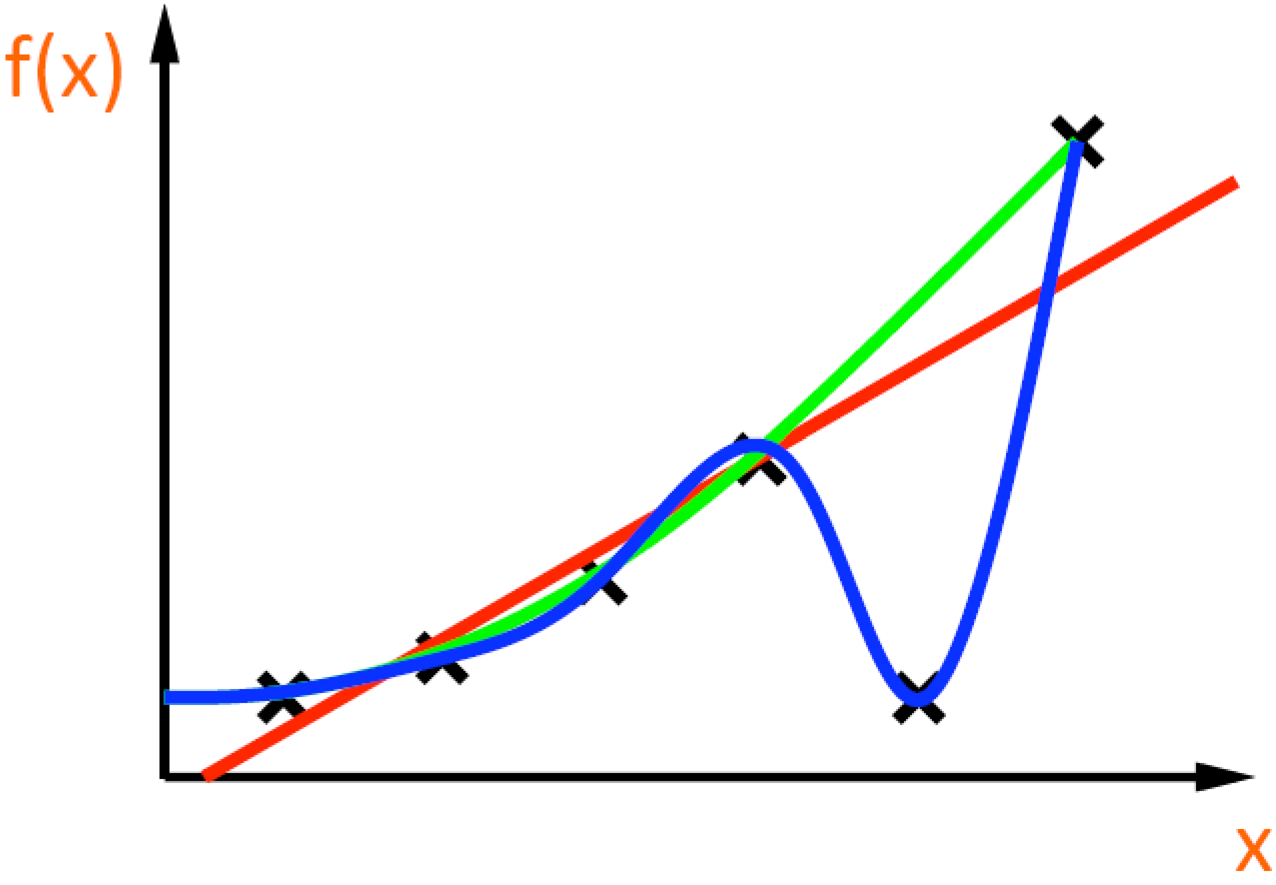

Are complicated models preferred?

- Order = 9?

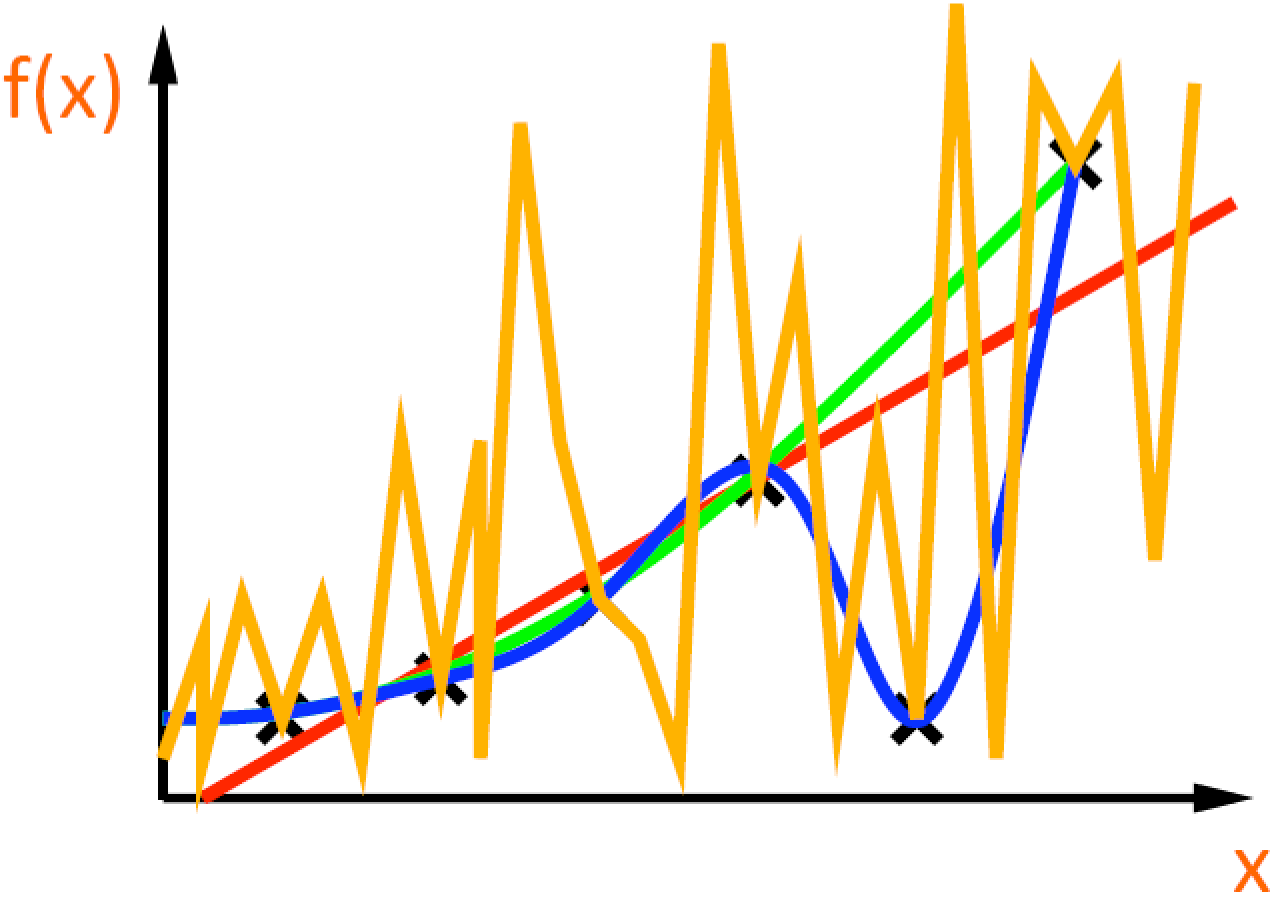

Are complicated models preferred?

Advice #2: use ML & sophisticated models when necessary

- Issues with model selection!!

Unsupervised Learning

- Task: learn an explanatory function $f(x), \; x \in \mathcal{X}$
- Aka "Learning without a teacher"

Feature space $\mathcal{X}$

Words in documents

⇨

Word distribution
(probability of a word)

- No training/test split

Unsupervised Learning – density estimation

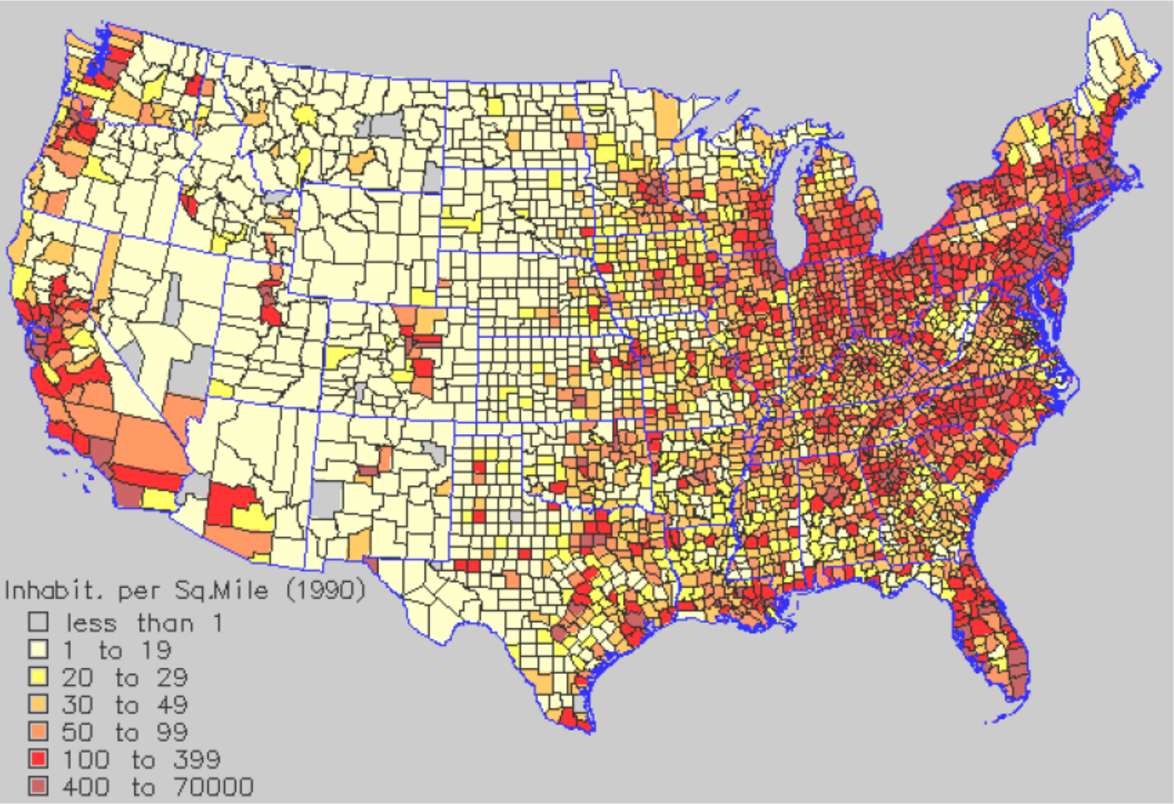

Feature space $\mathcal{X}$
geographical information of a location

Density function
$f(x), \ x \in \mathcal{X}$

Unsupervised Learning – clustering

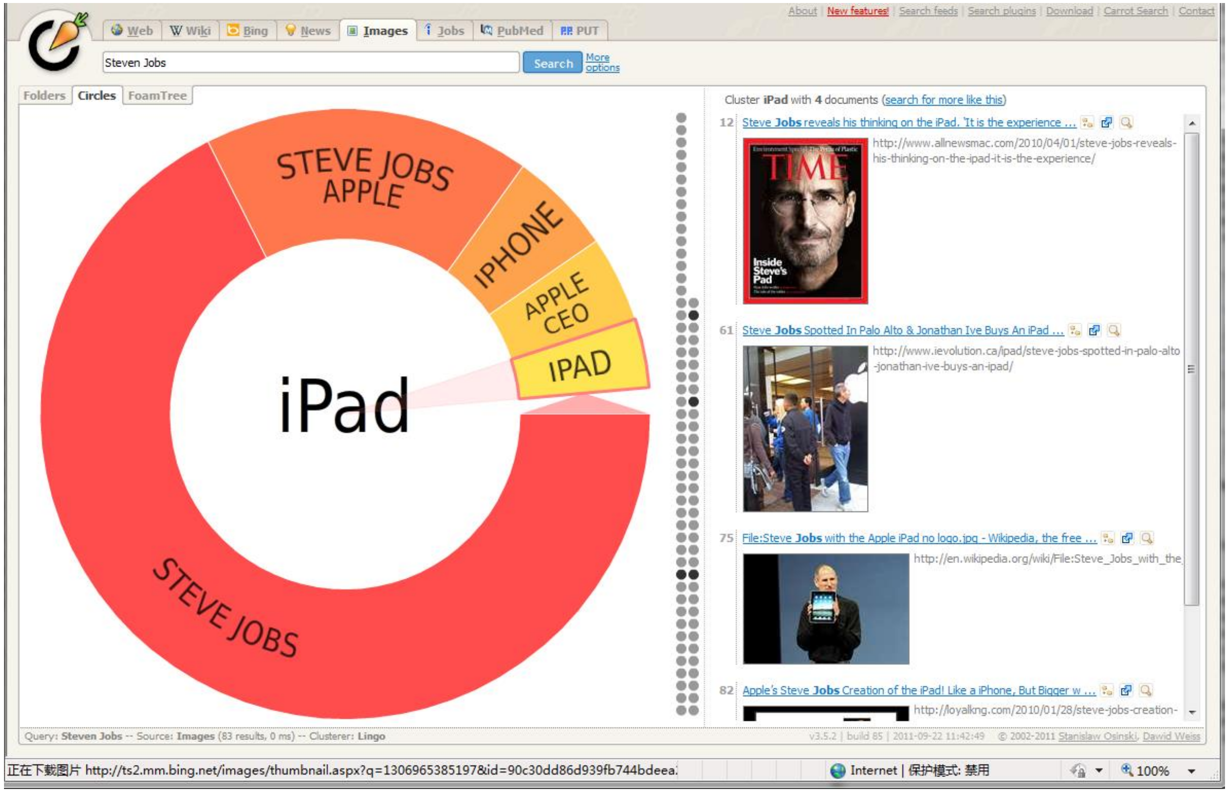

Cluster **iPad** with **4** documents (search for more like this)

12 Steve **Jobs** reveals his thinking on the iPad. 'It is the experience ...
http://www.allnewsmac.com/2010/04/01/steve-jobs-reveals-his-thinking-on-the-ipad-it-is-the-experience/

61 Steve **Jobs** Spotted In Palo Alto & Jonathan Ive Buys An iPad ...
http://www.ievolution.ca/ipad/steve-jobs-spotted-in-palo-alto-jonathan-ive-buys-an-ipad/

75 File:Steve **Jobs** with the Apple iPad no logo.jpg - Wikipedia, the free ...
http://en.wikipedia.org/wiki/File:Steve_Jobs_with_the

82 Apple's Steve **Jobs** Creation of the iPad! Like a iPhone, But Bigger w ...
http://loyalkng.com/2010/01/28/steve-jobs-creation-

Query: **Steven Jobs** -- Source: **Images** (83 results, 0 ms) -- Clusterer: **Lingo**

v3.5.2 | build 85 | 2011-09-22 11:42:49 © 2002-2011 Stanislaw Osinski, Dawid Weiss

正在下载图片 http://ts2.mm.bing.net/images/thumbnail.aspx?q=1306965385197&id=90c30dd86d939fb744bdeea:

http://search.carrot2.org/stable/search

Feature space $\mathcal{X}$
Attributes (e.g., pixels & text) of images

Cluster assignment function
$f(x), \ x \in \mathcal{X}$

Deep Generative Models

- Learn a generative model

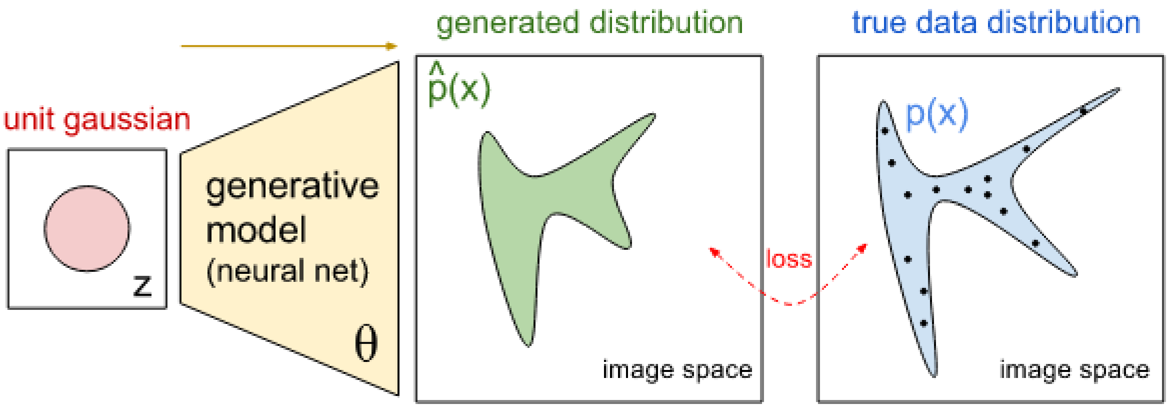

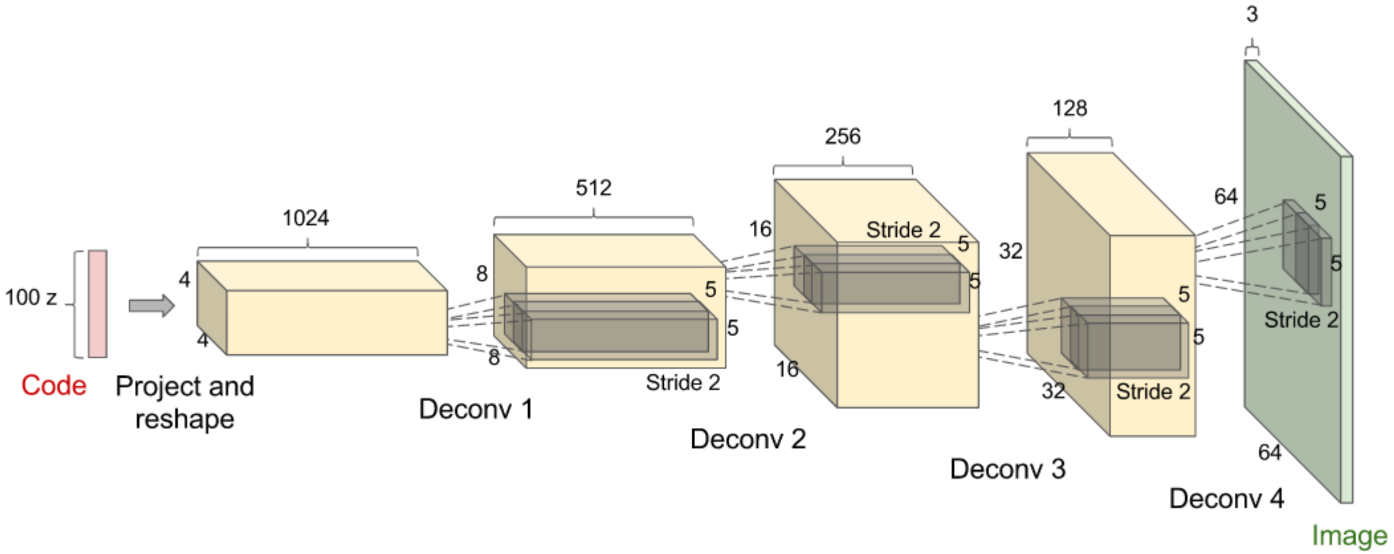

Unsupervised Learning – dimensionality reduction

Images have thousands or millions of pixels

Can we give each image a coordinate, such that similar images are near each other ?

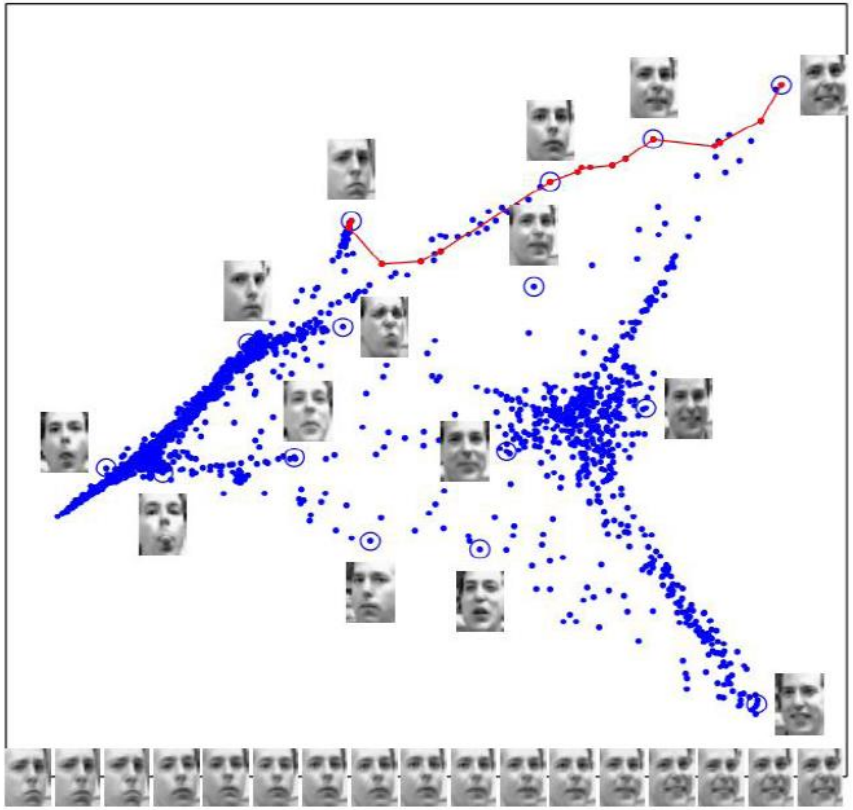

Feature space $\mathcal{X}$
pixels of images

Coordinate function in 2D space
$f(x), \ x \in \mathcal{X}$

Reinforcement Learning

- Critical Component of Intelligence
  - Understanding and advancing how an artificial agent can learn to make good decisions to do new tasks is fundamental challenge in artificial intelligence and machine learning

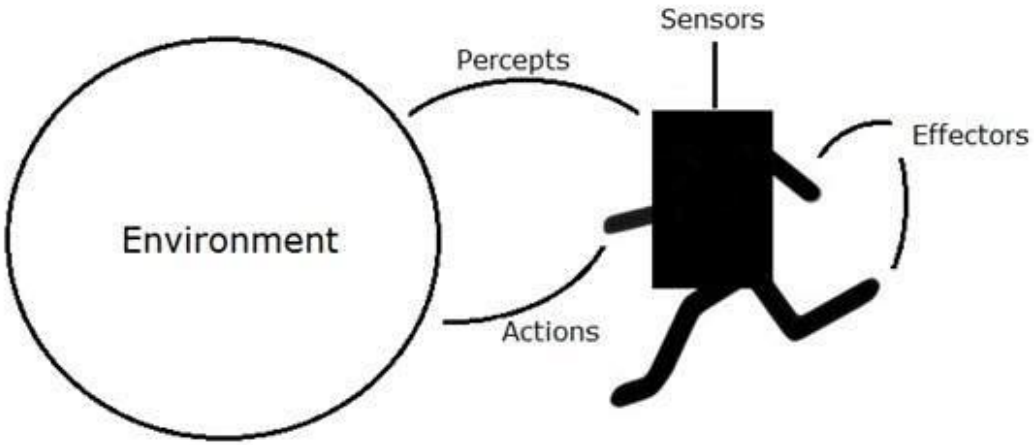

Reinforcement Learning (RL)

- Learn to make good sequences of decisions
  - Sequences of decisions:
    - repeated interactions with the world
  - Good:
    - reward for sequence of decisions
  - Learn:
    - don’t know in advance how world works

Reinforcement Learning (RL)

- RL deals with agents that must sense & act upon their environment.

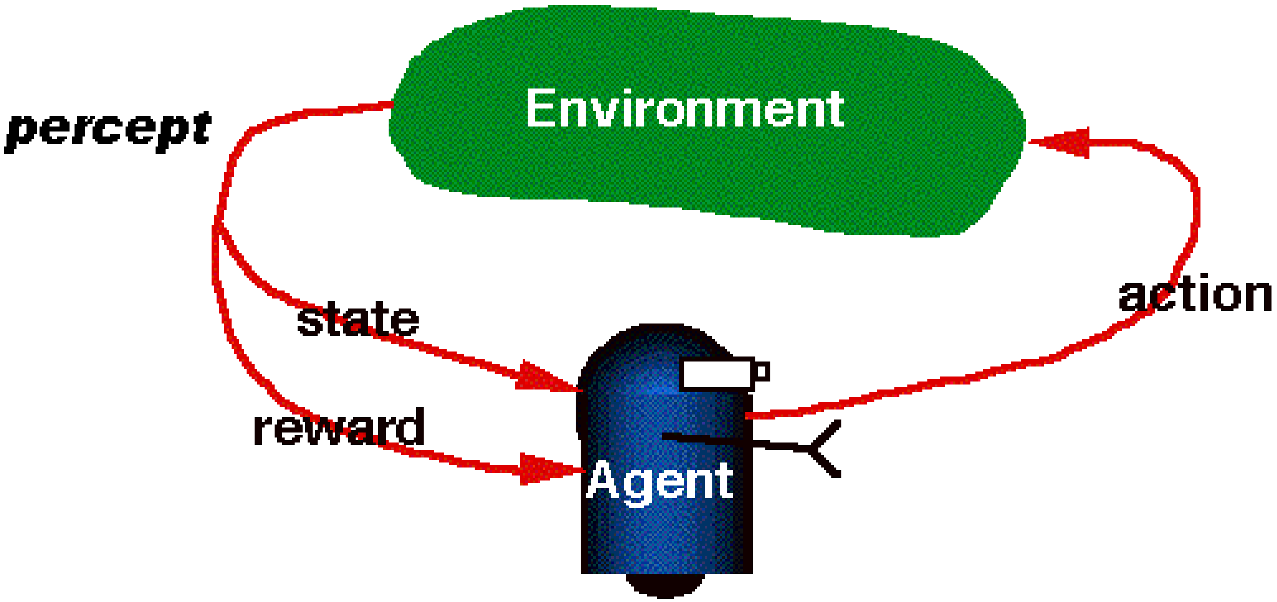

- This combines classical AI and machine learning techniques.
- It the most comprehensive problem setting.

Summary: what is machine learning

- Machine Learning seeks to develop theories and computer systems for

dealing with

- complex, real world data, based on the system's own experience with data, and (hopefully) under a unified model or mathematical framework, that

have nice properties.

Summary: what is machine learning

- Machine Learning seeks to develop theories and computer systems for
  - representing;
  - classifying, clustering, recognizing, organizing;
  - reasoning under uncertainty;
  - predicting;
  - and reacting to
  - …
- complex, real world data, based on the system's own experience with data, and (hopefully) under a unified model or mathematical framework, that

have nice properties.

Summary: what is machine learning

- Machine Learning seeks to develop theories and computer systems for
  - representing;
  - classifying, clustering, recognizing, organizing;
  - reasoning under uncertainty;
  - predicting;
  - and reacting to
  - …
- complex, real world data, based on the system's own experience with data, and (hopefully) under a unified model or mathematical framework, that
  - can be formally characterized and analyzed;
  - can take into account human prior knowledge;
  - can generalize and adapt across data and domains;
  - can operate automatically and autonomously;
  - and can be interpreted and perceived by human.

- ML covers algorithms, theory and very exciting applications
- It’s going to be fun and challenging ☺

Resources for Further Learning

- Top-tier Conferences:
  - International Conference on Machine Learning (ICML)
  - Advances in Neural Information Processing Systems (NIPS)
  - Uncertainty in Artificial Intelligence (UAI)
  - International Joint Conference on Artificial Intelligence (IJCAI)
  - AAAI Annual Conference (AAAI)
  - Artificial Intelligence and Statistics (AISTATS)

- Top-tier Journals:
  - Journal of Machine Learning Research (JMLR)
  - Machine Learning (MLJ)
  - IEEE Trans. on Pattern Recognition and Machine Intelligence (PAMI)
  - Artificial Intelligence
  - Journal of Artificial Intelligence Research (JAIR)
  - Neural Computation

Hot Topics from ICML & NIPS

- Hot topics:
  - Deep Learning with Rich Model Architecture
  - Probabilistic Latent Variable Models & Bayesian Nonparametrics
  - Sparse Learning in High Dimensions
  - Large-scale Optimization and Inference
  - Online learning
  - Reinforcement Learning
  - Learning Theory
  - Interdisciplinary Research on Machine Learning, Cognitive Science , etc.

Resources for Further Learning

- Text books:
  - Pattern Recognition and Machine Learning
  - Probabilistic Graphical Models (http://pgm.stanford.edu/)
  - Deep Learning

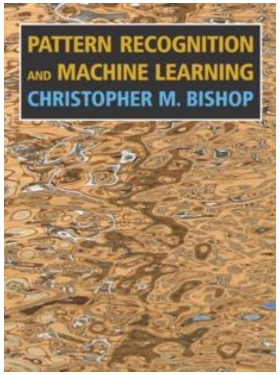

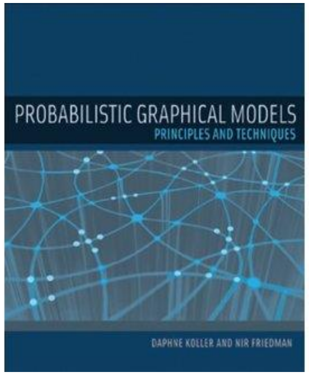
PROBABILISTIC GRAPHICAL MODELS
PRINCIPLES AND TECHNIQUES
DAPHNE KOLLER AND NIR FRIEDMAN

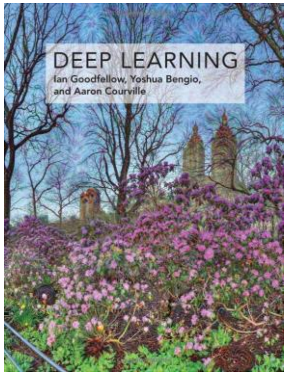

- Public lectures:
  - CMU :
    - http://www.cs.cmu.edu/~guestrin/Class/10708-F08/projects.html
  - Stanford:
    - http://cs228.stanford.edu/
    - http://cs228t.stanford.edu/
  - UPenn:
    - http://www.seas.upenn.edu/~cis620/

Thanks!

Jun Zhu

dcszj@mail.tsinghua.edu.cn

http://ml.cs.tsinghua.edu.cn/~jun