

[70240413 Statistical Machine Learning, Spring, 2018]

Statistical Machine Learning

Theory and Applications

Jun Zhu

`dcszj@mail.tsinghua.edu.cn`

`http://ml.cs.Tsinghua.edu.cn/~jun`

State Key Lab of Intelligent Technology & Systems
Tsinghua University

February 27, 2018

A bit about the Instructor

- ◆ Jun Zhu, Associate Professor, Depart. of Computer Science. I received Ph.D. in 2009. My research interest includes machine learning, Bayesian methods, and data mining
- ◆ I did post-doc at the Machine Learning Department in CMU with Prof. Eric P. Xing. Before that I was invited to visit CMU for twice. I was also invited to visit Stanford for joint research (with Prof. Li Fei-Fei)
- ◆ 2015: Adjunct Associate Professor at CMU
- ◆ Published 100+ papers on the top-tier ML conferences and journals, including JMLR, TPAMI, ICML, NIPS, etc.
- ◆ Served as Area Chairs for ICML, NIPS, UAI, AAAI, IJCAI; Associate Editor for PAMI, AI Journal
- ◆ Research is supported by National 973, NSFC, “Tsinghua 221 Basic Research Plan for Young Talents”.
- ◆ IEEE AI’s 10 to Watch; MIT TR35 China (pioneers)
- ◆ Homepage: <http://ml.cs.tsinghua.edu.cn/~jun>



Contact Information

◆ Jun Zhu

- State Key Lab of Intelligent Technology and Systems,
Department of Computer Science, Tsinghua U.

- Office: Rm 4-513, FIT Building
- E-mail: dcszj@tsinghua.edu.cn
- Phone: 62772322, 18810502646
- Office hours: Thursday afternoon 3:00pm-5:00pm
 - Better to make an appointment in advance

Teaching Assistants

◆ Jiaxin Shi (Head TA)

- E-mail: ishijiaxin@126.com
- Phone: 62795869, 18810690095
- Bayesian methods, Deep learning
- Publish at VAST, NIPS, ICLR



◆ Yucen Luo (Head TA)

- E-mail: luoyucencen@163.com
- Phone: 62795869, 18810301080
- Deep learning, Latent variable models
- Publish at ICML, CVPR.



Teaching Assistants

◆ Jialian Li

- E-mail: lijialian7@163.com
- Phone: 18510243737
- Reinforcement learning

◆ Kun Xu

- E-mail: vofhqn@gmail.com
- Phone: 15701006589
- Bayesian methods, Bandits

◆ Yichi Zhou



- E-mail: vofhqn@gmail.com
- Phone: 15701006589
- Bayesian methods, Bandits

◆ Haosheng Zou

- E-mail: zouhs16@mails.tsinghua.edu.cn
- Phone: 18800120568
- Reinforcement learning

◆ TA office hours: [Wed afternoon 3:00pm-5:00pm](#)

◆ Office: [Rm 1-508/509, FIT Building; 62795869](#)



Resources

◆ Mainly class slides/notes

◆ Recommended text books

- Christopher M. Bishop. *Pattern Recognition and Machine Learning*, Springer, 2007.
- Trevor Hastie, Robert Tibshirani, Jerome Friedman. *Elements of Statistical Learning*. 2nd Edition, Springer, 2009.

◆ Further readings:

- Conferences:
 - Theory: ICML, NIPS, UAI, COLT, AISTATS, AAAI, IJCAI
 - App: KDD, SIGIR, WWW, ACL
- Journals:
 - JMLR, PAMI, MLJ

Prerequisites

- ◆ Knowledge of probability, linear algebra, statistics and algorithms
 - Calculus:
 - Derivative, integral of multivariate functions
 - Linear Algebra
 - Matrix inversion, eigen-decomposition, ...
 - Basic Probability and Statistics
 - Probability distributions, Mean, Variance, Conditional probabilities, Bayes rule, ...

- ◆ Knowledge of programming languages, e.g., C/C++, Java, matlab, Python

- ◆ **Homework 0:** take the Self-Evaluation
 - Minimum & modest background tests (available at course webpage)

Overview of Class

- ◆ Introduction
- ◆ Unsupervised learning
- ◆ Supervised learning
- ◆ Learning theory
- ◆ Probabilistic graphical models
- ◆ Bayesian methods
- ◆ Sparse learning
- ◆ Deep learning
- ◆ Reinforcement Learning

3 units	
6 units	HW1 out
6 units	
3 units	HW1 due HW2 out
6 units	
3 units	HW2 due HW3 out
3 units	
6 units	HW3 due HW4 out
6 units	
	HW4 due June 7

Grading

◆ Participation (10%)

- 1 mid-term quiz (10 points each time)


◆ Homeworks (40%)

- 4 homeworks (10 points each time)

◆ Project (50%)

- 2~4 students to form a team
- Apply machine learning to solve a real problem
 - Choose one task at Kaggle (<http://www.kaggle.com/competitions>)
- Submit materials:
 - a proposal (6th week), a mid-term report (9th week), a final report (18th week), and the implementation code (18th week)
- All reports should be in NIPS format, written in English:
(<http://nips.cc/Conferences/2014/PaperInformation/StyleFiles>)
- Poster presentation (16th or 17th week)

Some example Kaggle tasks




2018 Data Science Bowl

Find the nuclei in divergent images to advance medical discovery

Featured · 2 months to go · 📁 biology

\$100,000
1,743 teams



Google Landmark Recognition Challenge

Label famous (and not-so-famous) landmarks in images

Research · 3 months to go · 📁 image data

\$2,500
81 teams



Digit Recognizer

Classify handwritten digits using the famous MNIST data

Getting Started · 3 years to go · 2,361 kernels

1,422 teams



Titanic: Machine Learning from Disaster

Predict survival on the Titanic using Excel, Python, R & Random Forests

Getting Started · 3 years to go · 6,074 kernels

5,864 teams

Other Projects

◆ Self-defined topics

- Need to propose as early as possible to filter out improper ones

◆ Other candidates

- Chinese handwritten characters generation and recognition
- Adversarial attacks and defense of deep learning
- Reinforcement learning
- More to come

- ◆ If the end date is later than June 12, report the position in the leaderboard;
- ◆ Otherwise, ask TAs to define a train/test split and compare your methods with 1 or 2 baselines.

About final report

◆ We expect to see

- Problems (**what?**)
- Motivations (**why?**)
- Techniques (**how?**)
- Results & Analysis (**did you verify what you claimed above?**)
- Conclusions

◆ The final report should look like a NIPS technical paper

Questions?